

Spark API

Tao Ruangyam, ING Analytics - Frankfurt Hub

Istanbul 2020



What will be covered in this session

- Intro to Spark API
- RDD
- DataFrame API
- Window functions
- Built-in functions
- UDFs

Introduction



col1	col2	col3
Row 1		
Row 2		
...		
Row N		



RDD
(Resilient Distributed Dataset)

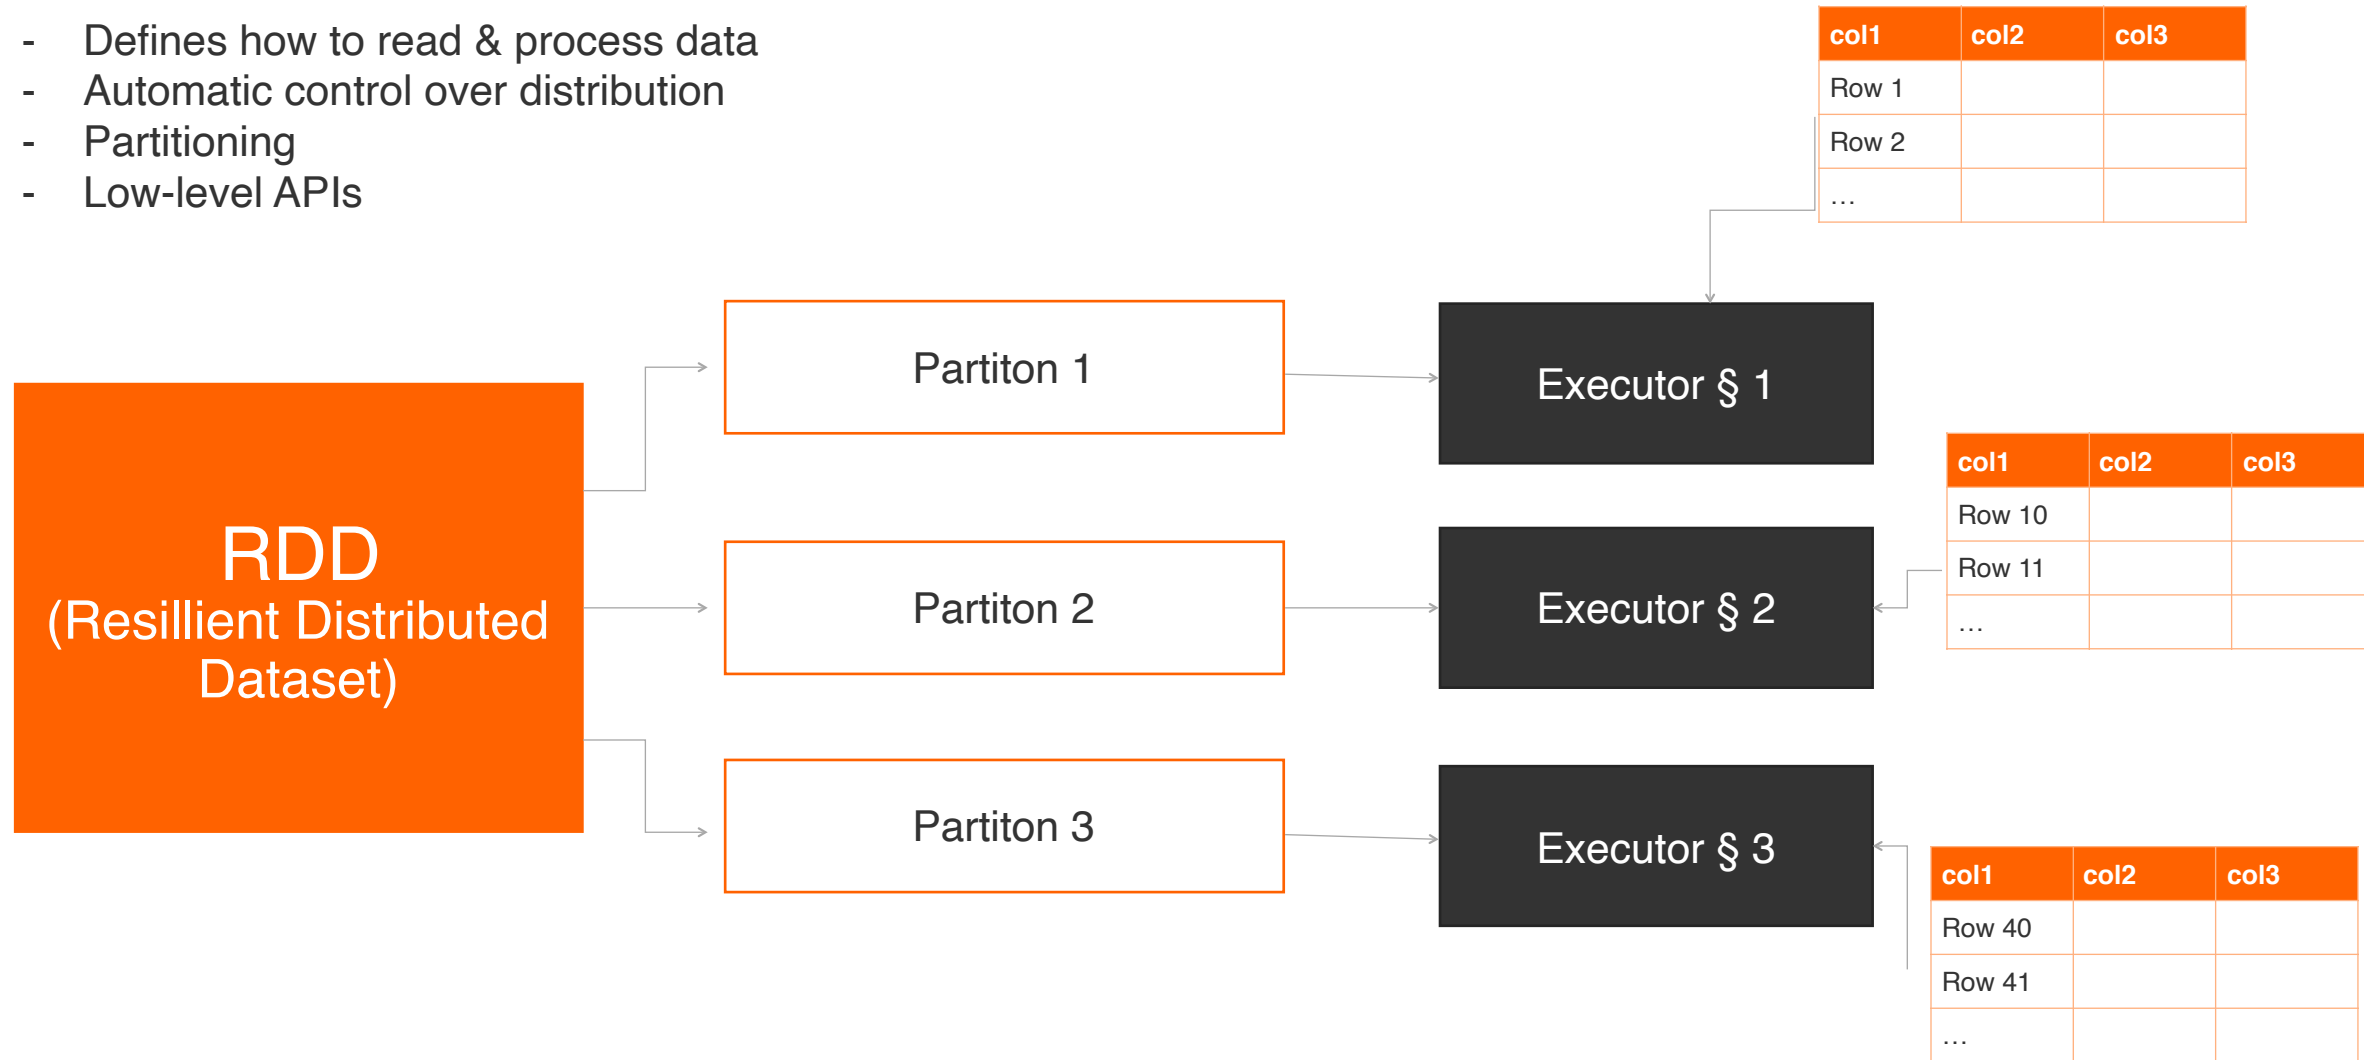


DataFrame
(abstraction over RDD)

With Serialisation


RDD

- Defines how to read & process data
- Automatic control over distribution
- Partitioning
- Low-level APIs



RDD

What coders see



API functions

- map()
- filter()
- cache()
- count()
- ...

RDD
(Resilient Distributed Dataset)

What it tells Driver

- **How the data is created**
 - From beginning
 - Until the output
- **How the data is distributed**
 - How to partition
 - Which executor has it

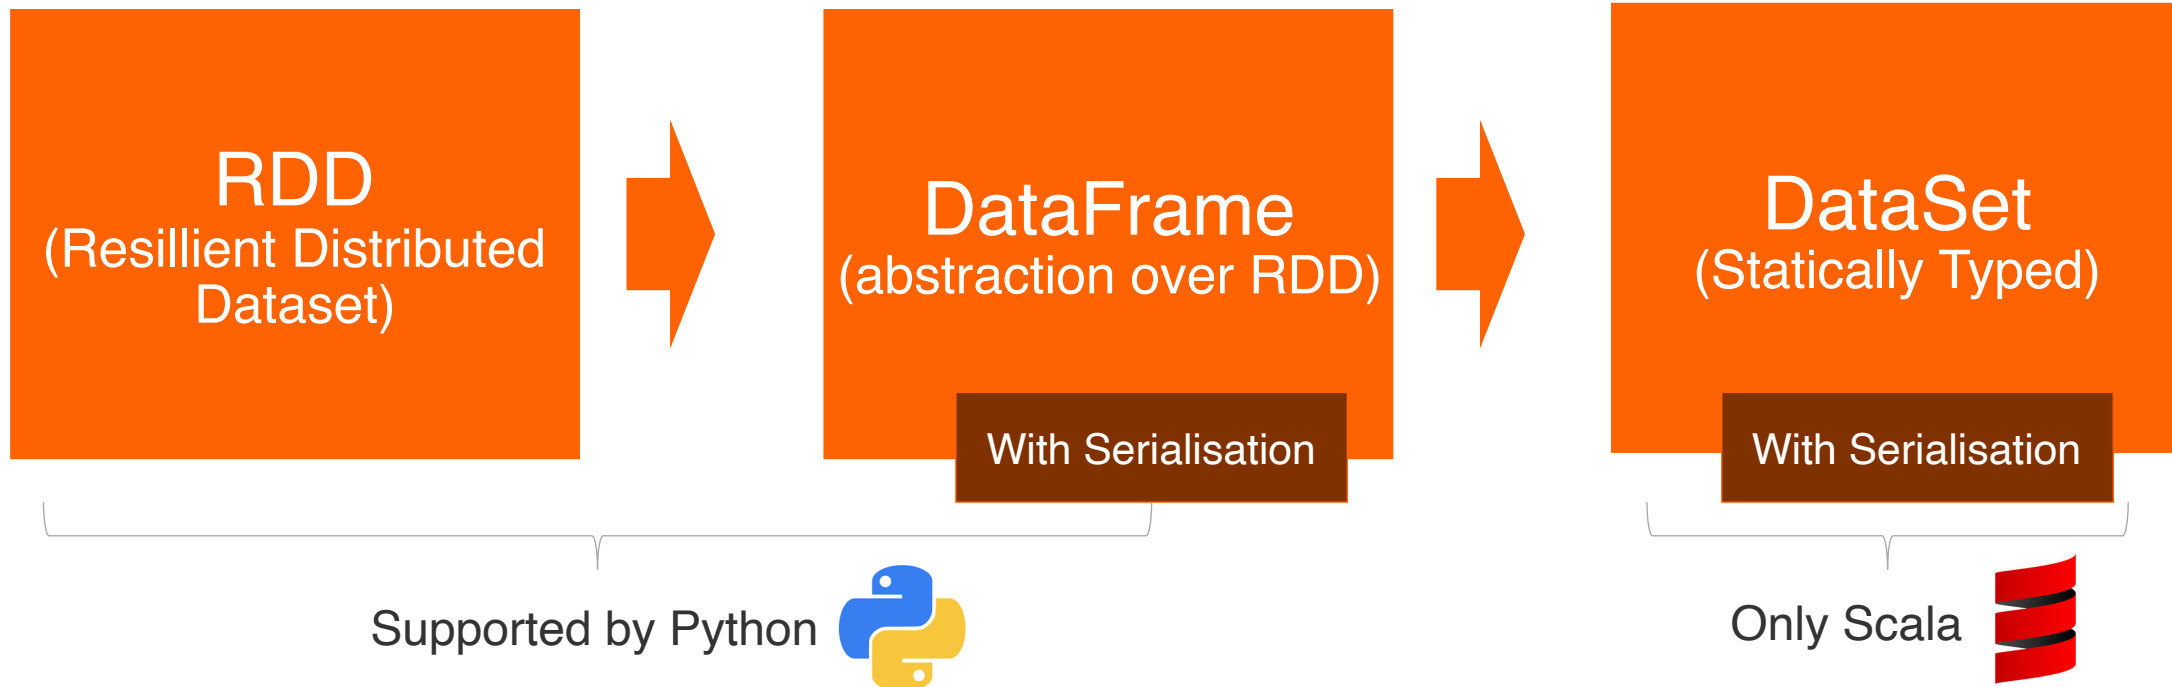
RDD

Create an RDD

```
rdd = spark.sparkContext.parallelize([
    ('Red', 22, 3500),
    ('Green', 70, 5500),
    ('Blue', 15, 9500)
])

rdd = spark.sparkContext \
    .textFile('/path/to/file.csv') \
    .map(lambda line: line.split(','))
```

API overview



DataFrame API

- Higher-level abstraction over **RDD**
- **No typesafe checking** at compilation (Dataset has)
- Fully compatible with **SparkSQL** API
- **Serialised!**

DataFrame API

Create a DataFrame

```
data = spark.read.csv('/path/to/file.csv')

data = spark.read.parquet('/path/to/file.parquet')

data = spark.sql('SELECT a,b,min(c) over (partition by a) FROM hivetable1')

data = spark.read \
    .format("kafka") \
    .option("kafka.bootstrap.servers", "localhost:9092") \
    .option("subscribe", "topic1") \
    .load()
```

DataFrame API

Filter, add columns, do some aggregation

```
from pyspark.sql.functions import *  
  
result = data.filter((col('a') < 10) | (col('a').isNull()))  
  
result = data.join(data2, "col-to-join", "inner")  
  
result = data.withColumn("newCol", explode("array-column"))
```

Note that, Spark is lazy evaluated

Code will not be executed right away

```
from pyspark.sql.functions import *
```

```
result = data \
    .filter((col('a') < 10) | (col('a').isNull())) \
    .join(data2, "col-to-join", "inner")
```

Not executed yet

```
exploded = result.withColumn("newCol", explode("array-column"))
```

Not executed yet

```
..  
..  
..  
..
```

```
exploded.show(25)
```

Executed!

```
n = exploded.count()
```

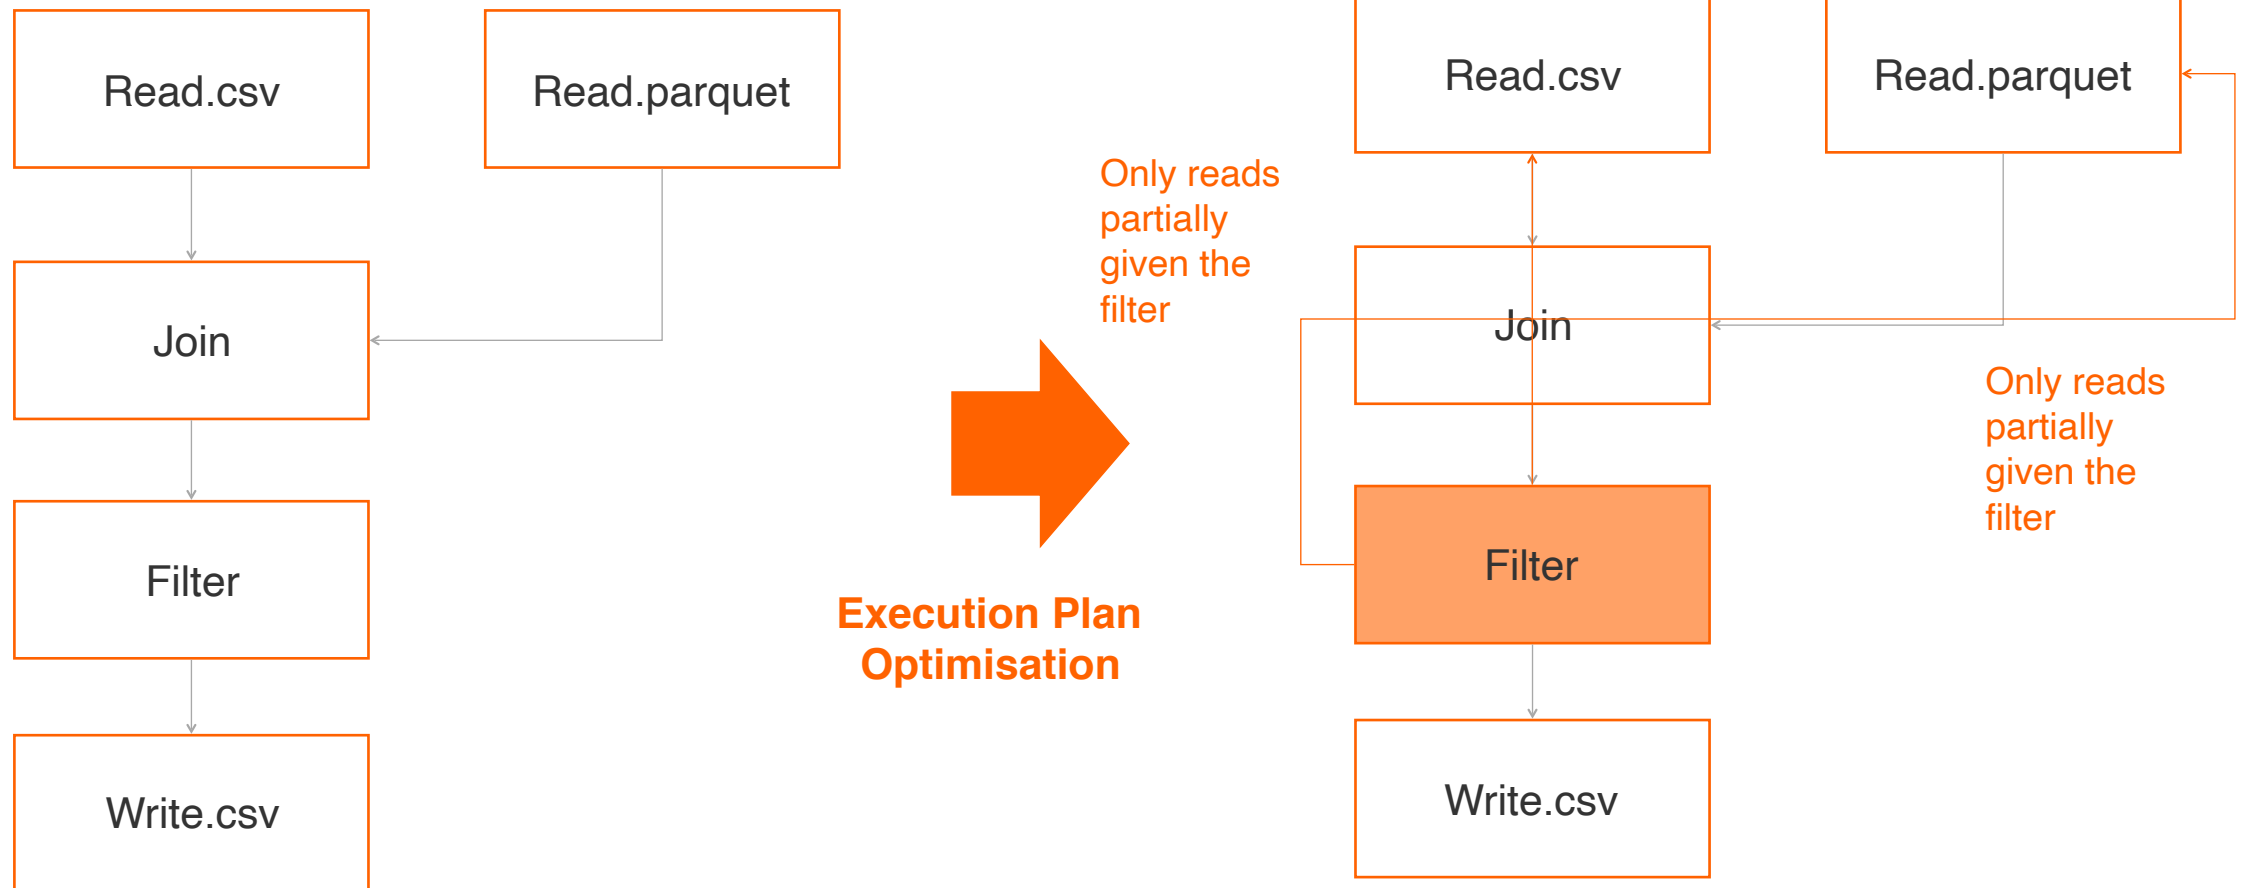
Executed!

```
exploded.write.save('path/to/file.parquet')
```

Executed!

Why lazy? Why not actively evaluate?

Process can be more optimisable with lazy evaluation

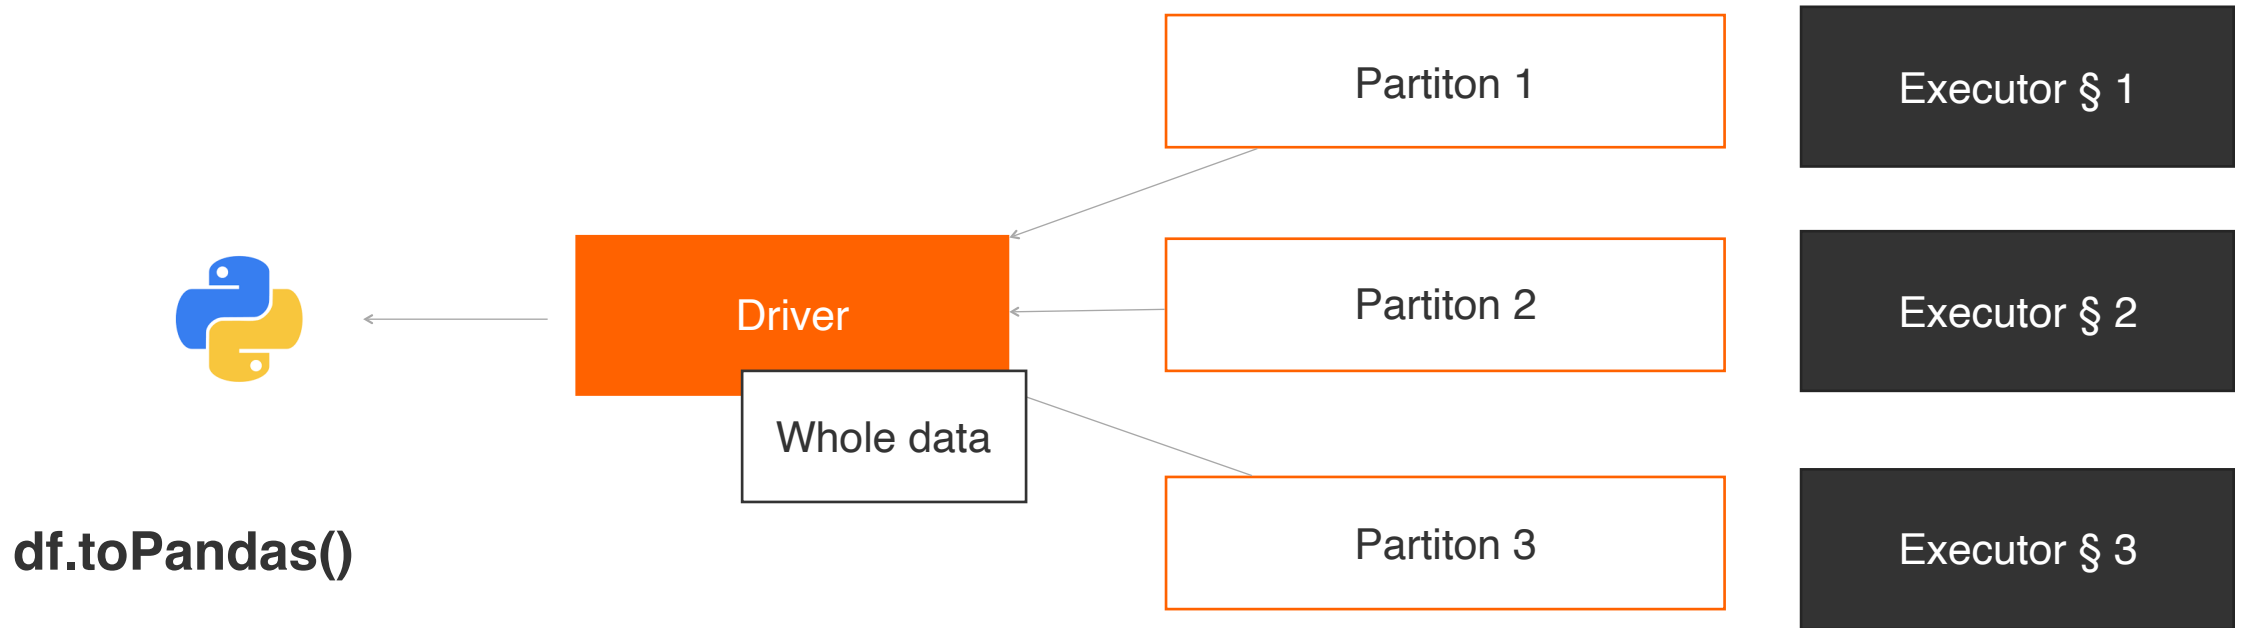


RDD API is accessible from DataFrame

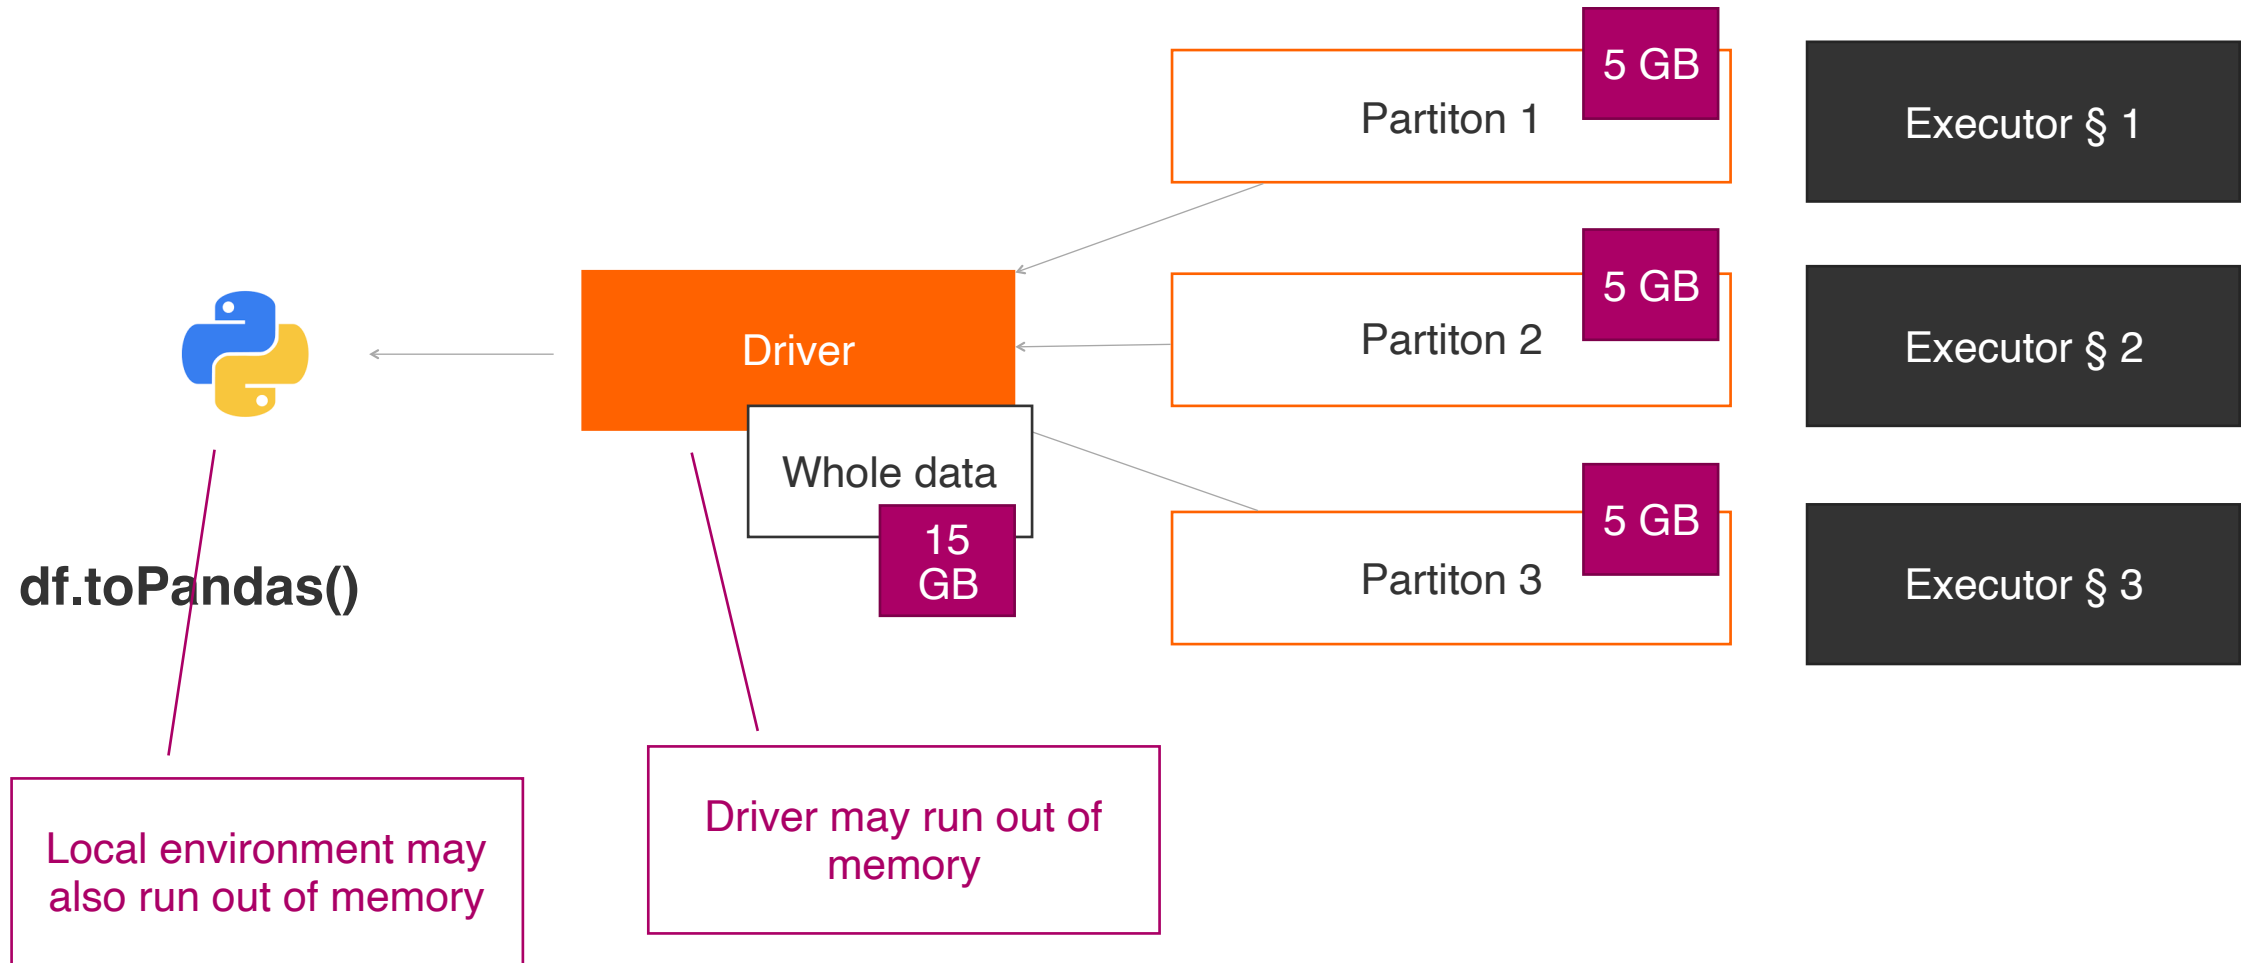
```
df = spark.read.parquet('path/to/file.parquet')  
out = df.rdd.map(lambda row: "Features: {}".format(row.feature))  
kv = df.rdd.keyBy(lambda row: row.a)
```

Since DataFrame is just an abstraction over RDD,
DF.rdd exposes the RDD inside

Ways to transfer data to Pandas



But this can cause out-of-memory




Pandas to Spark

```
import numpy as np
import pandas as pd

# Enable Arrow-based columnar data transfers
spark.conf.set("spark.sql.execution.arrow.enabled", "true")

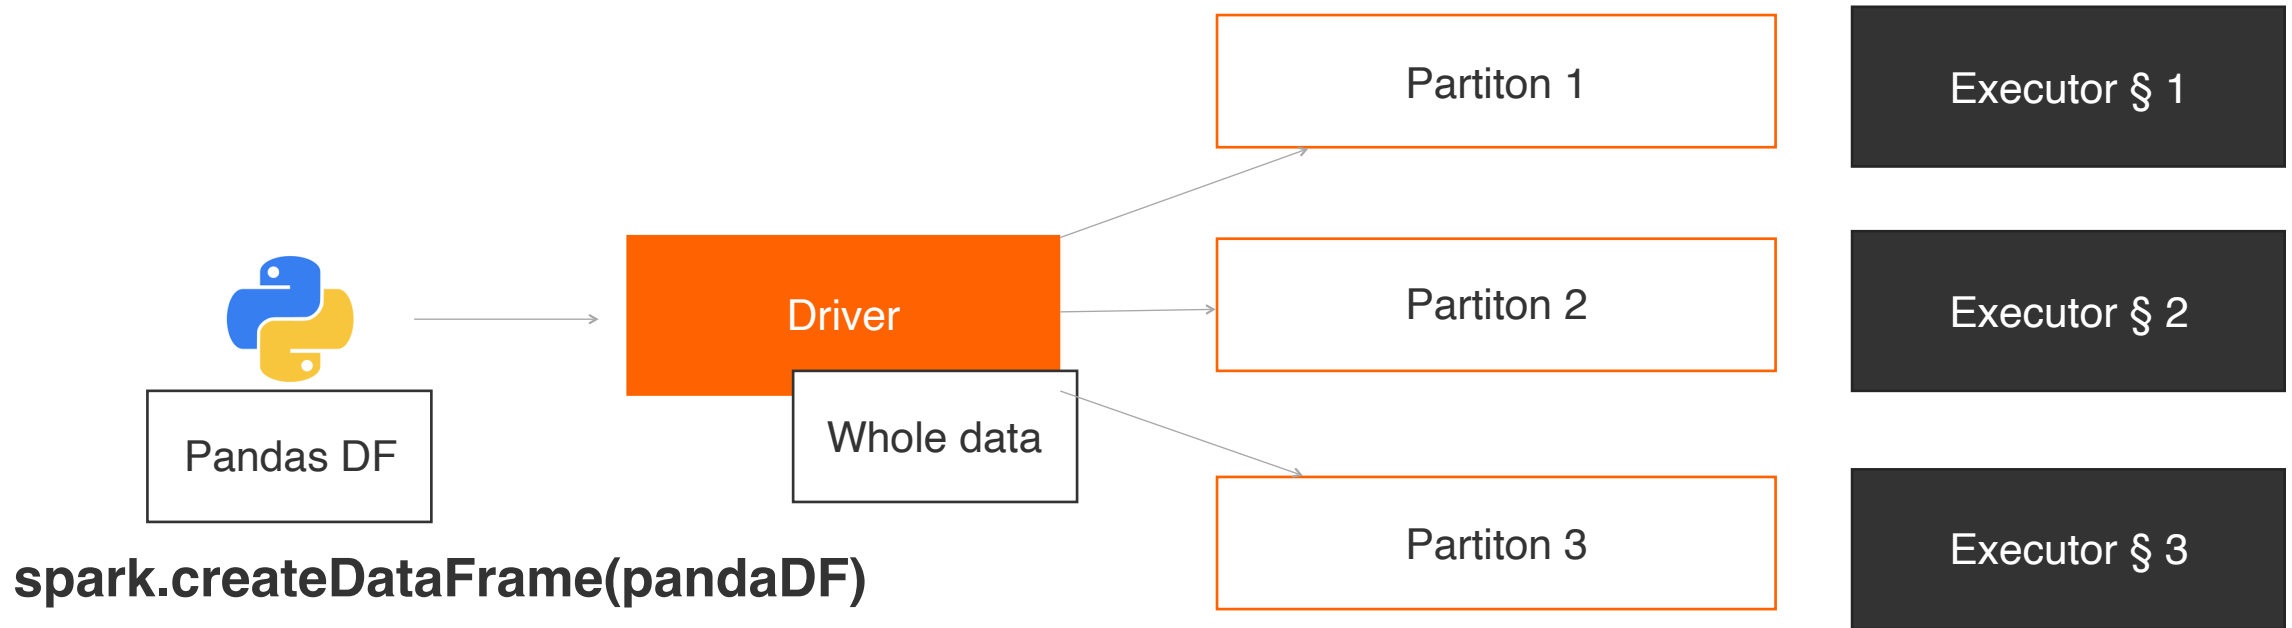
# Generate a pandas DataFrame
pdf = pd.DataFrame(np.random.rand(100, 3))

# Create a Spark DataFrame from a pandas DataFrame using Arrow
df = spark.createDataFrame(pdf)
```

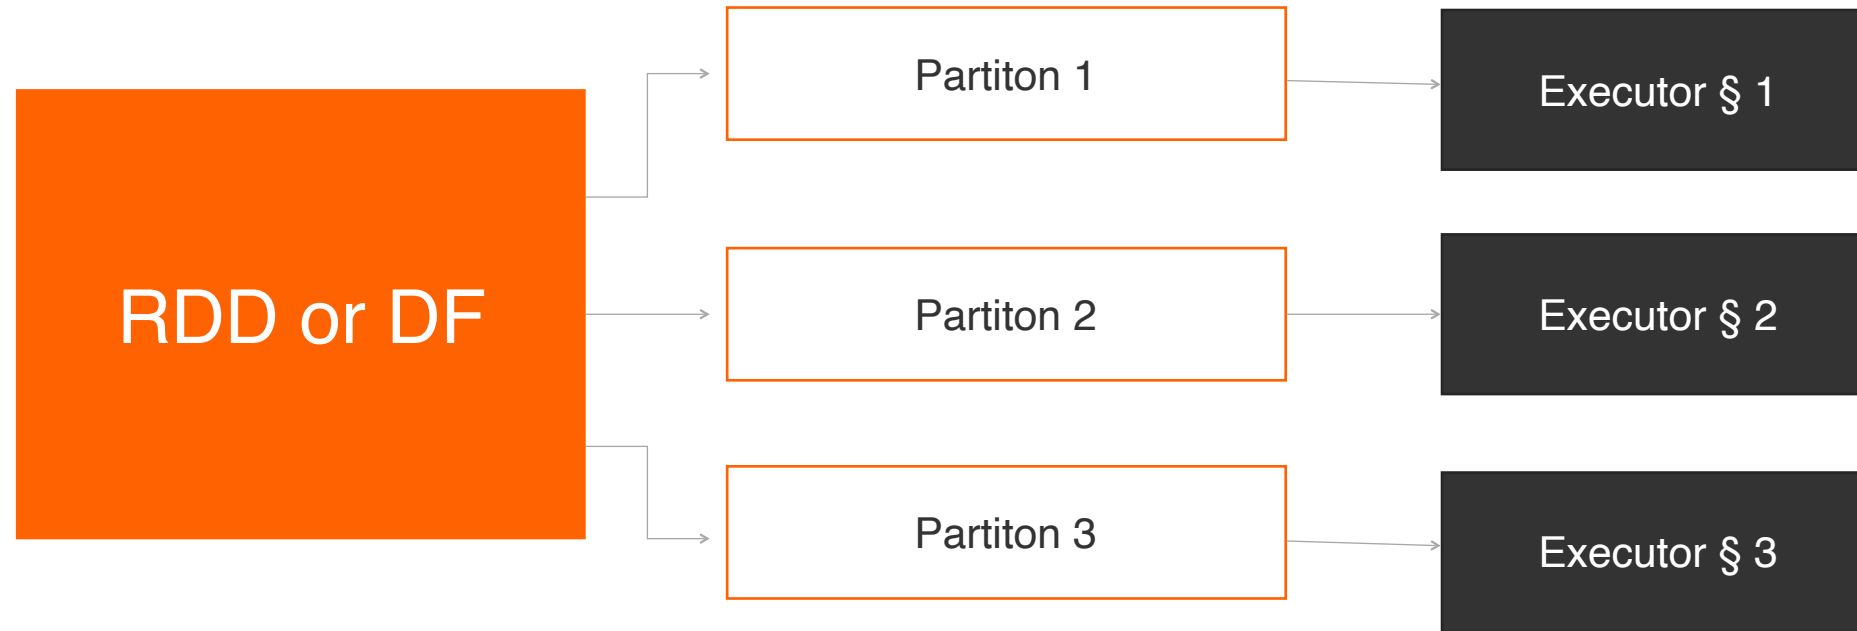


Arrow will transfer local Pandas dataframe to the Spark instance

Pandas to Spark

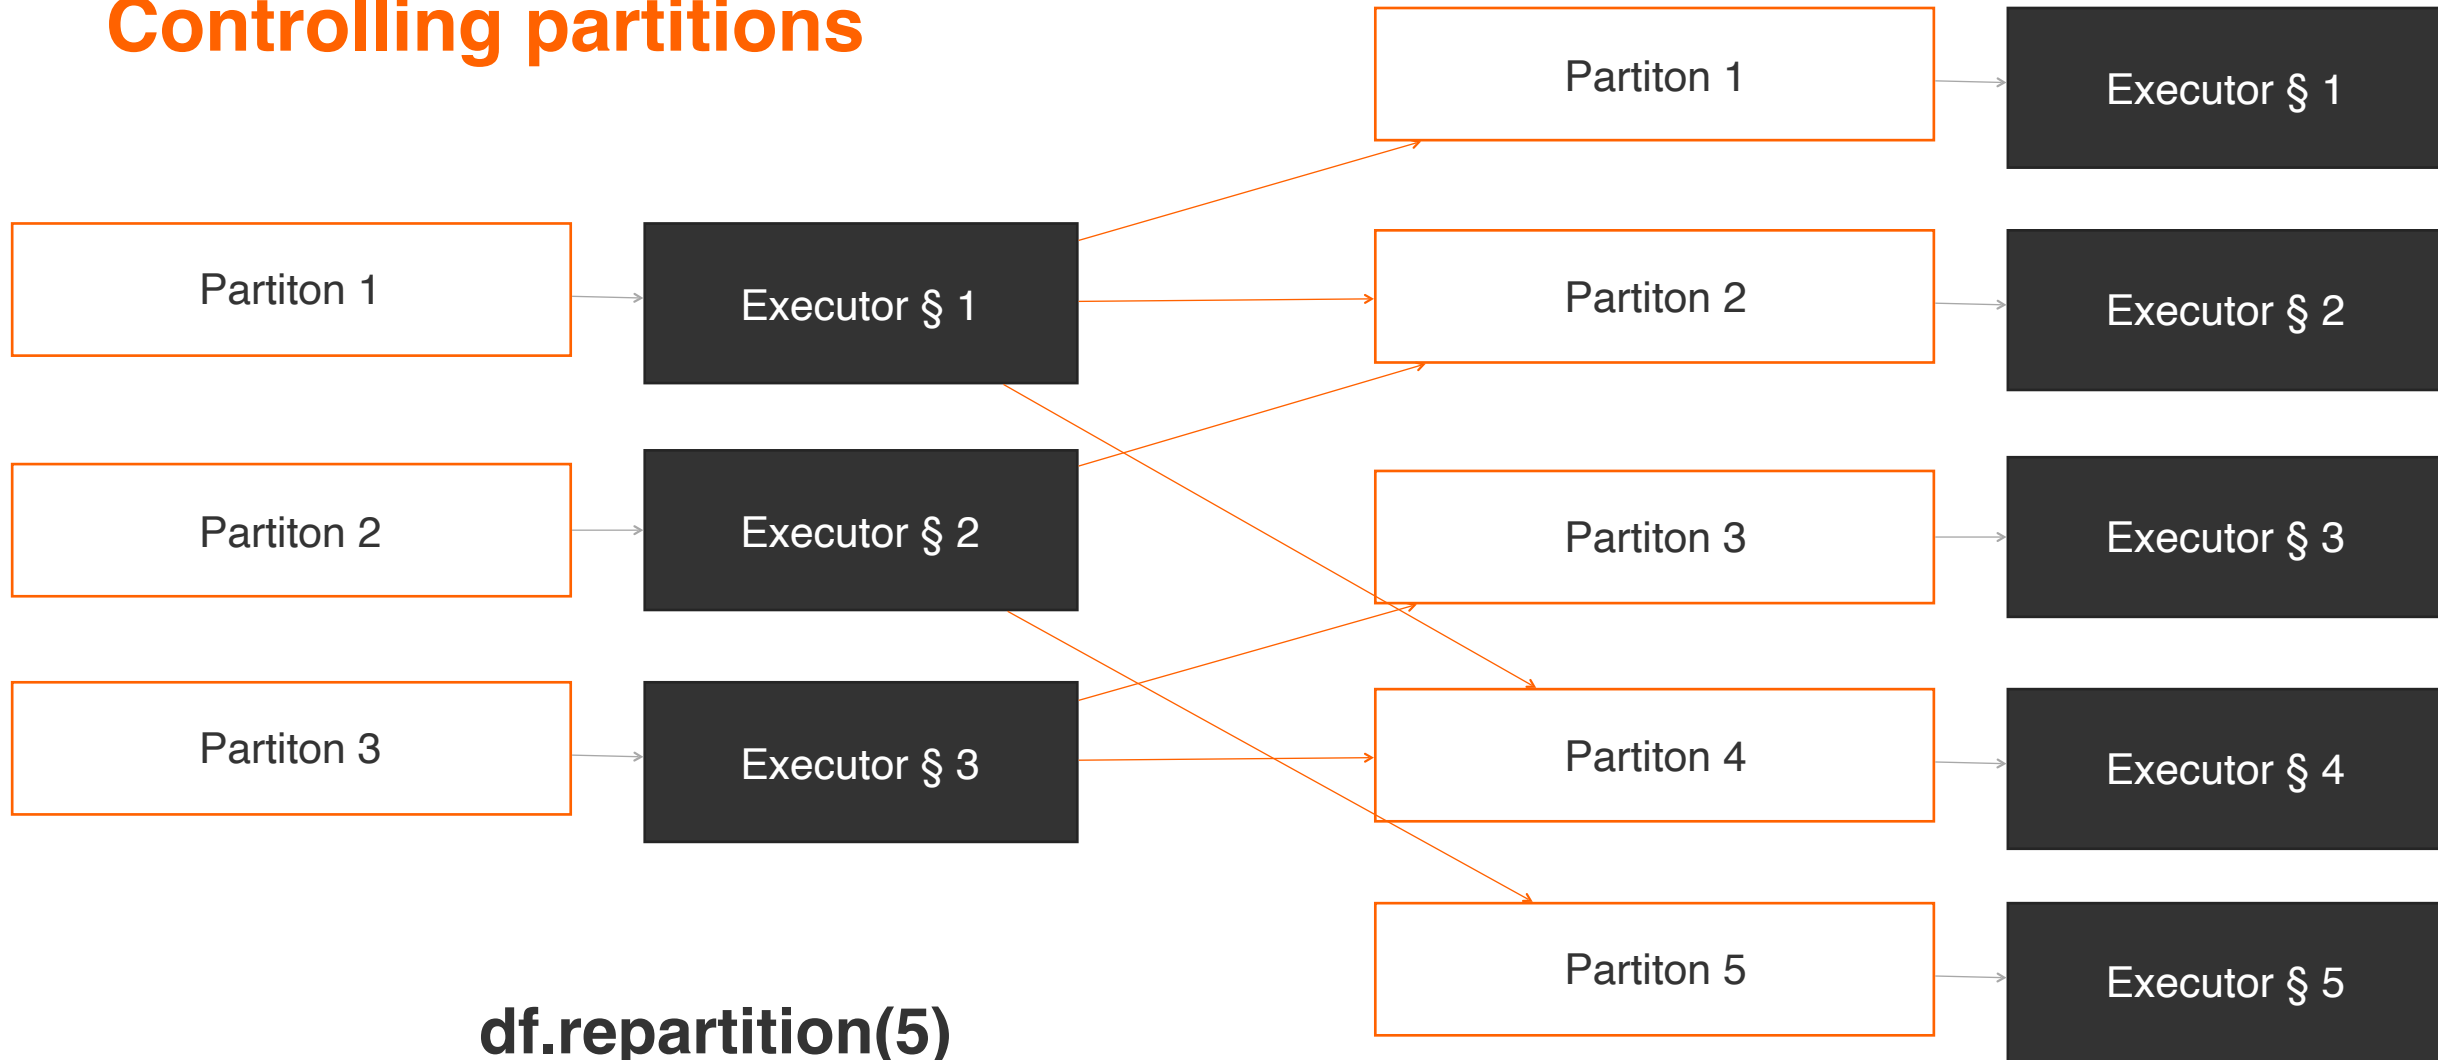


Controlling partitions



spark.default.parallelism (number of partitions after join or reduceByKey)

Controlling partitions

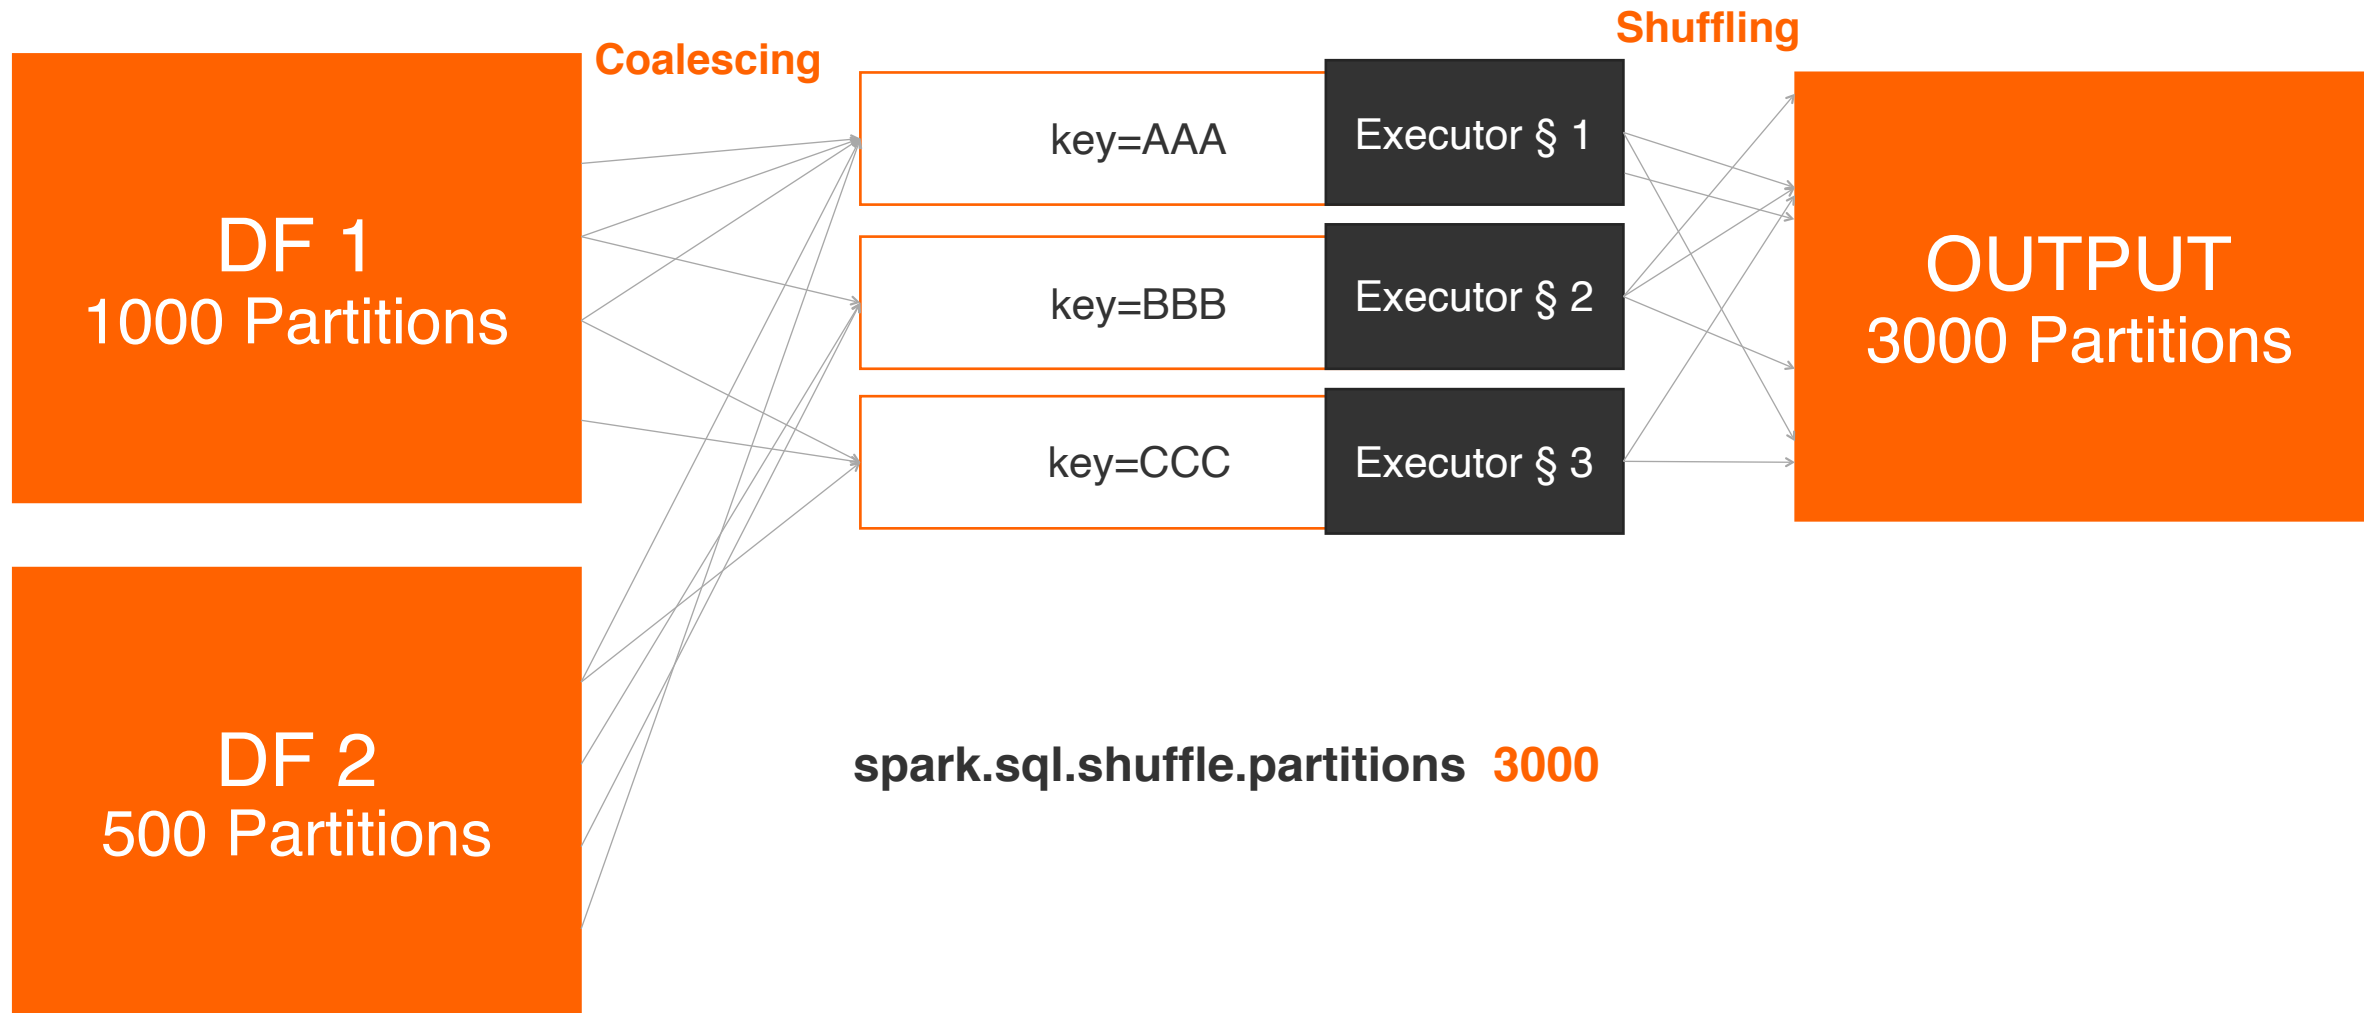


HashJoin

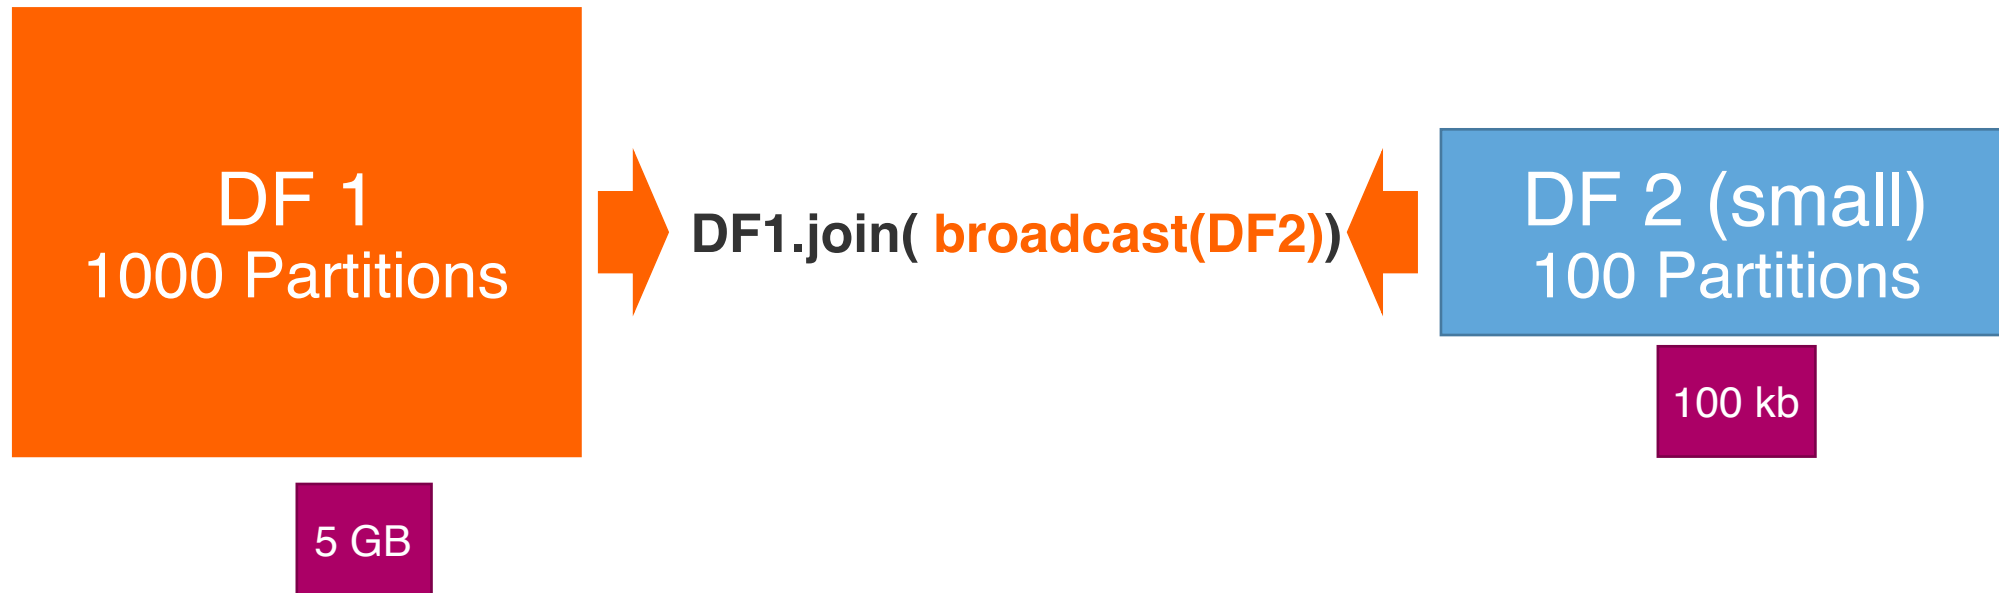


Regardless of prior number of partitions,
Spark **reshuffles** them by **joining keys**

HashJoin

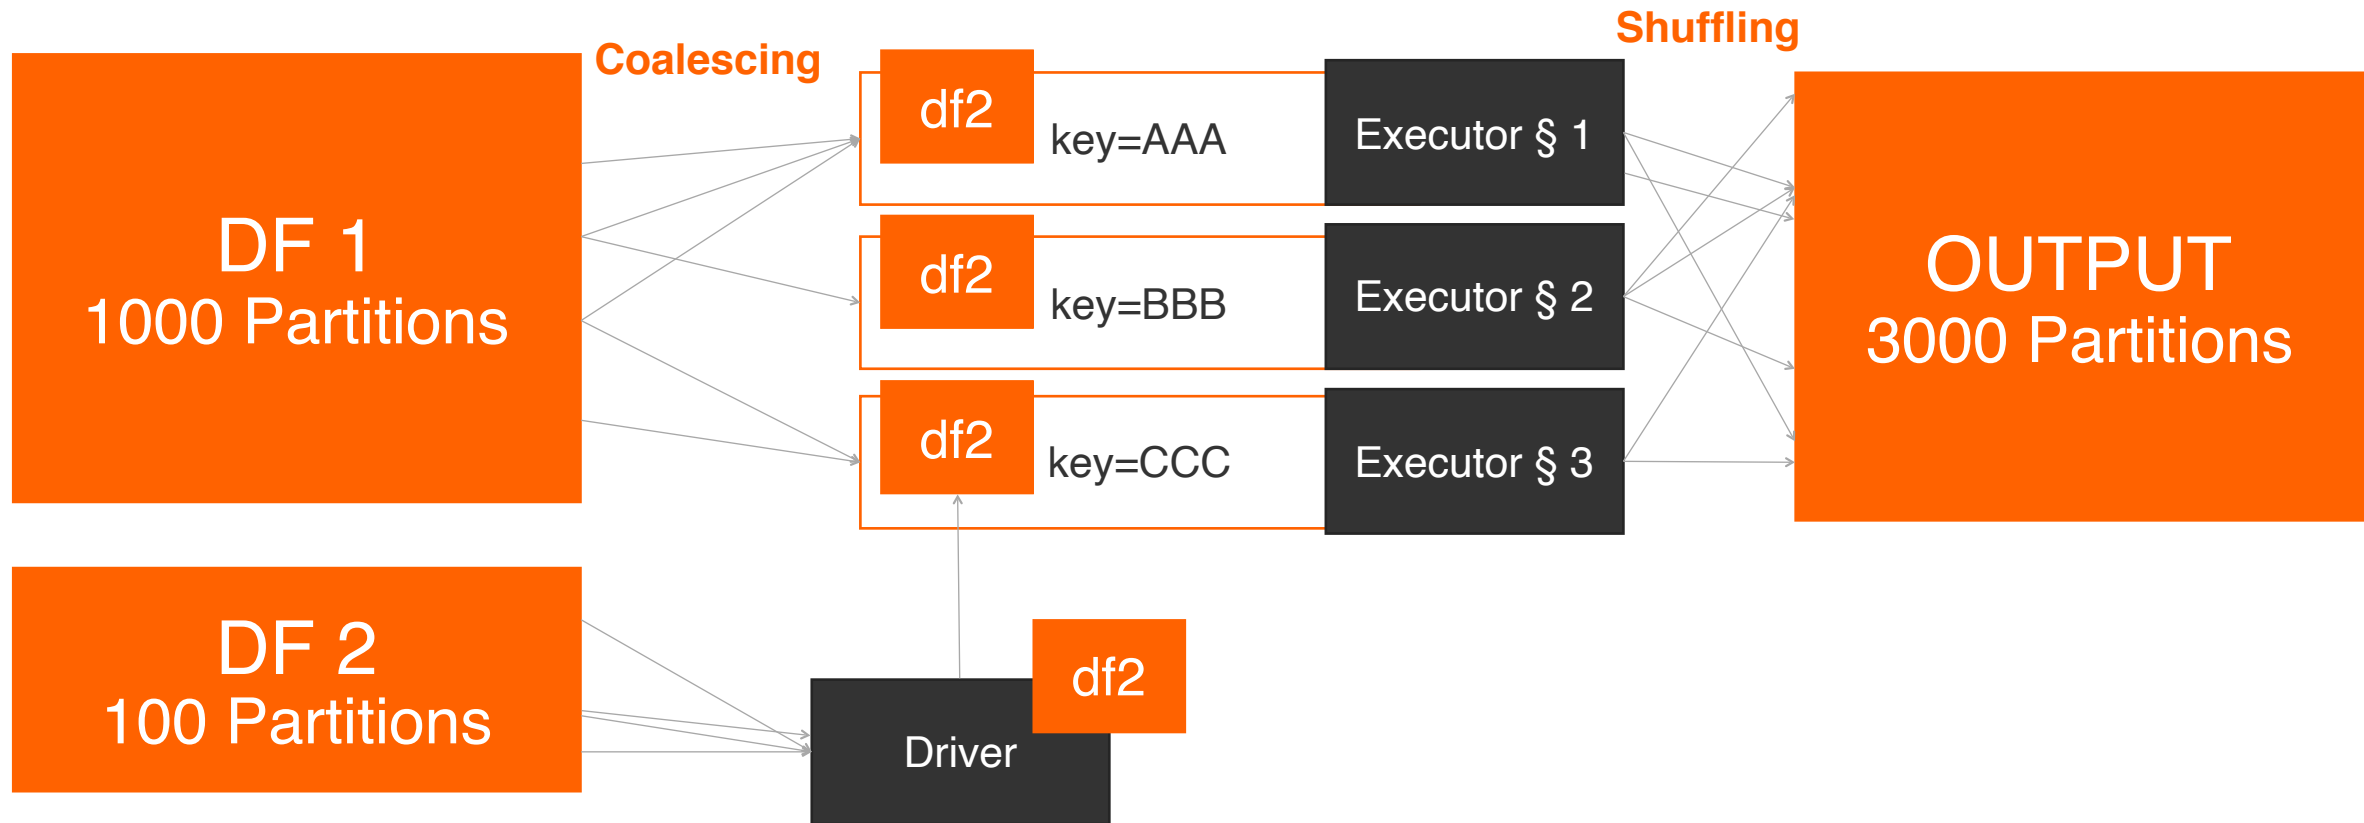


BroadcastHashJoin



Still works the same way with **HashJoin**
but all executors **get the whole broadcasted** dataframe

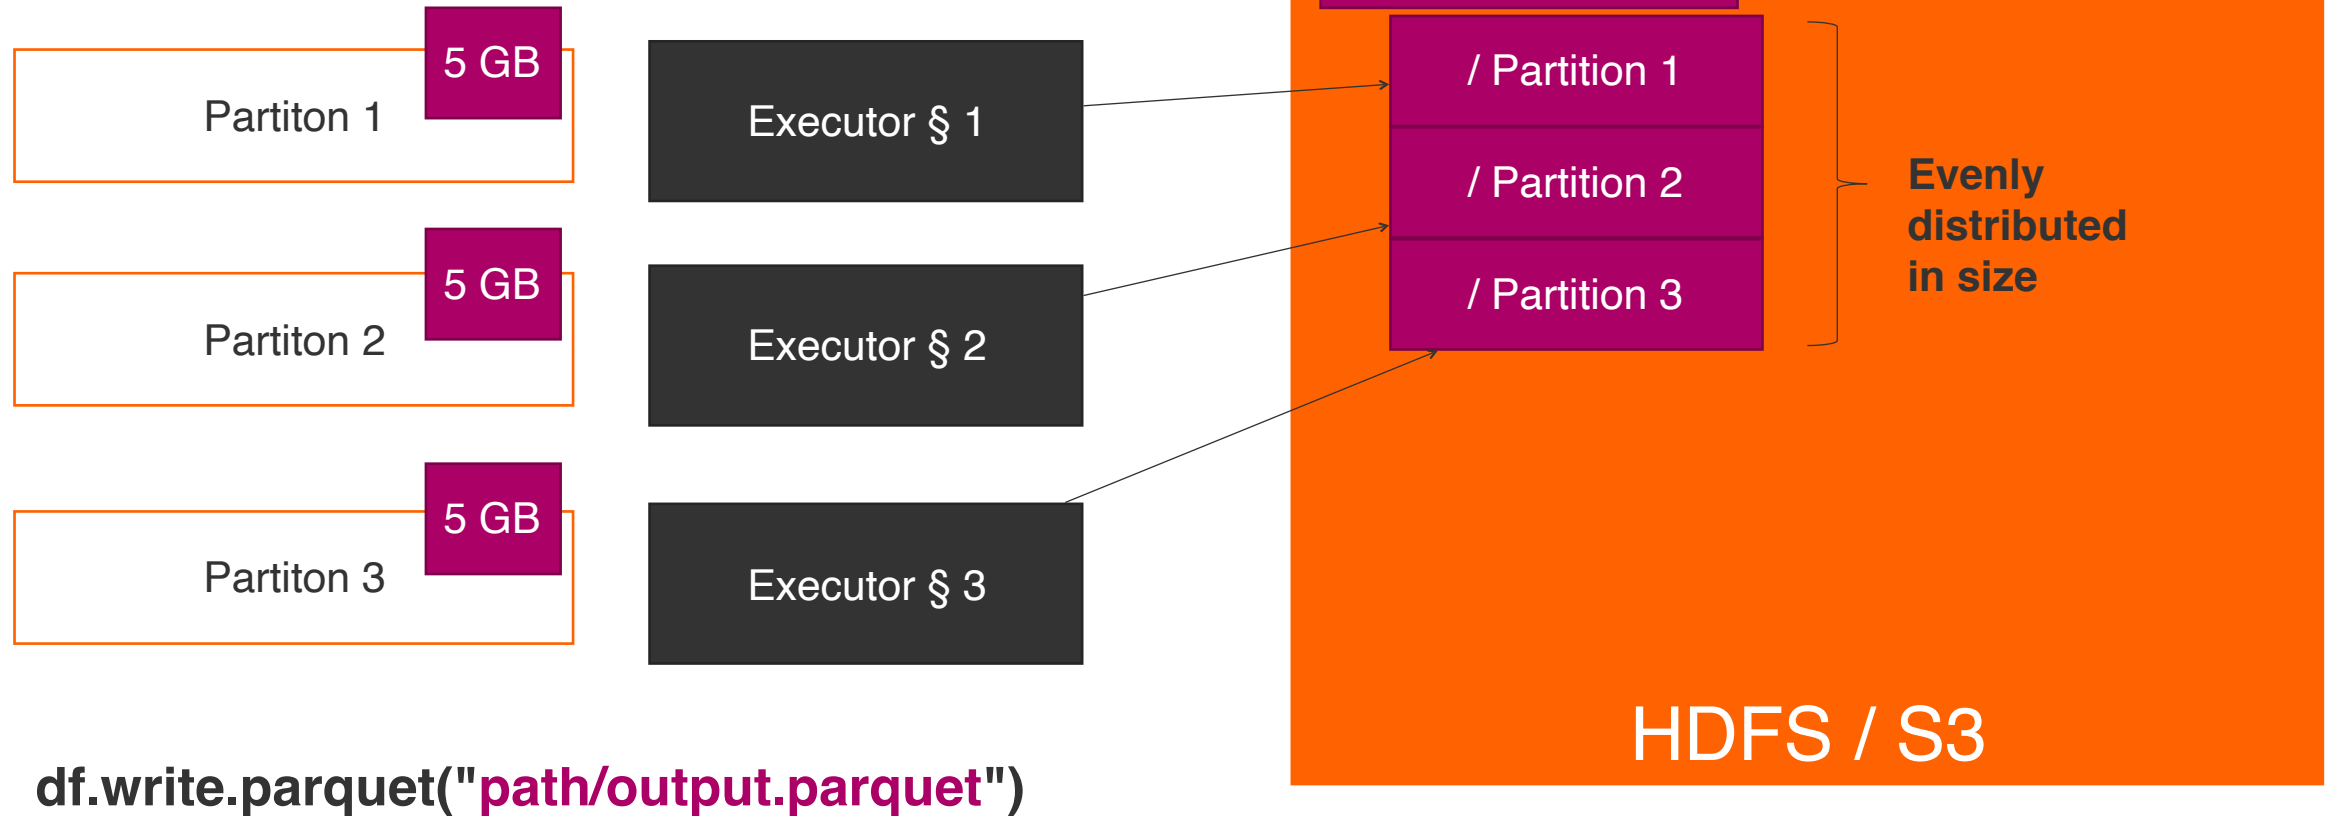
HashJoin



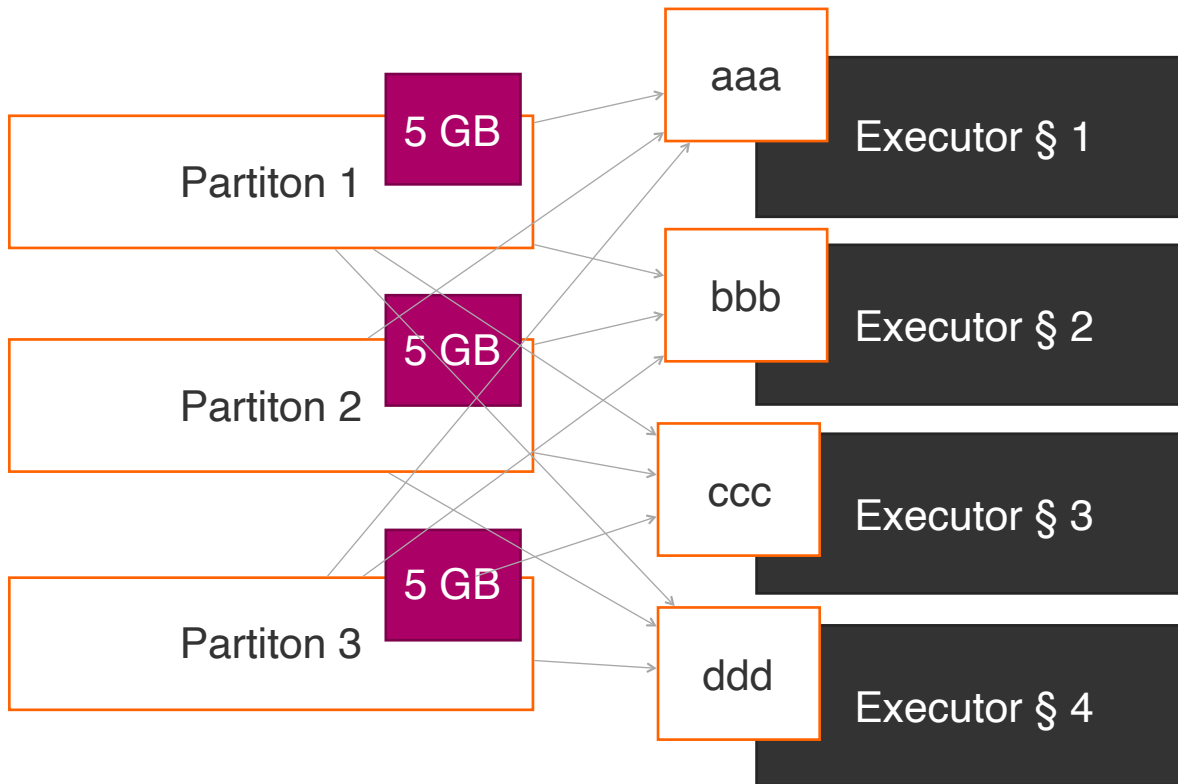
Writing output

```
df.write.option("header", "true").csv("path/file.csv")  
df.write.partitionBy("col1").parquet("path/file.parquet")  
df.write.mode("append").parquet("path/file.parquet")  
df.write.json("path/file.json")  
df.write.saveAsTable("schema.hive_table_name")
```

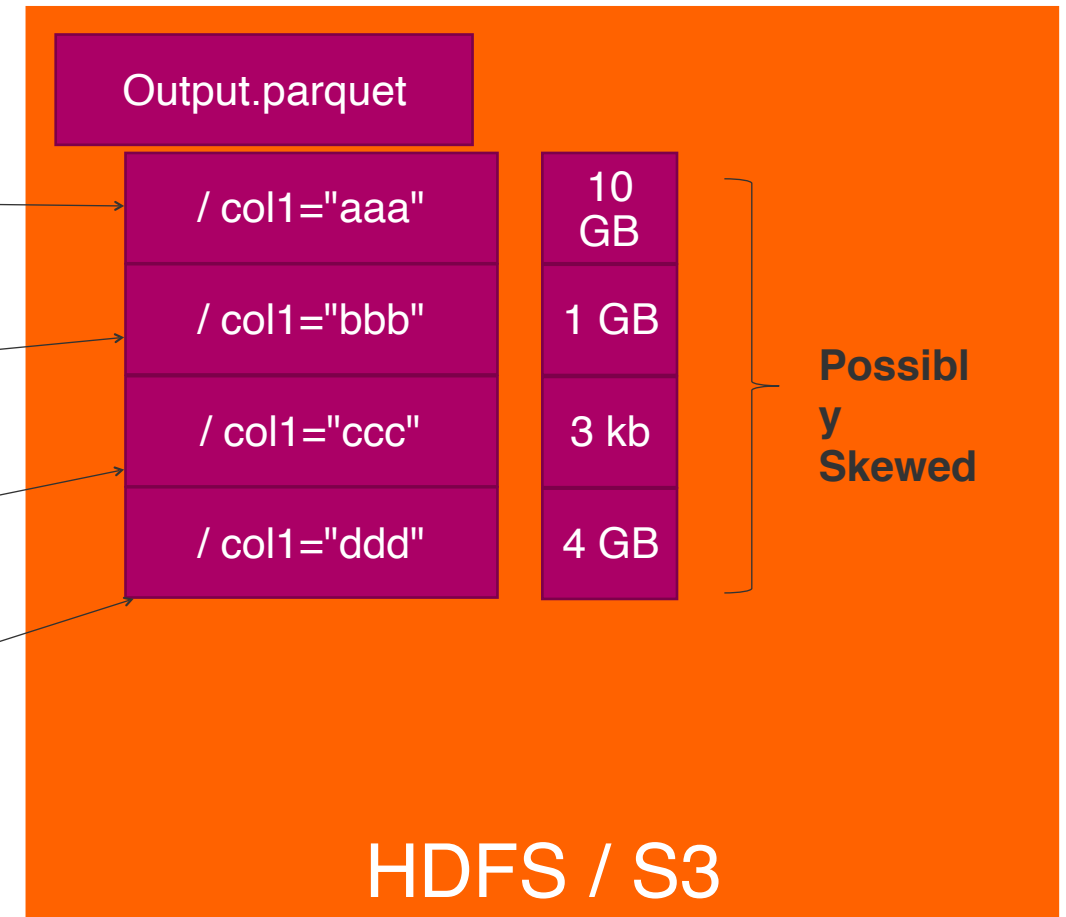

Writing output



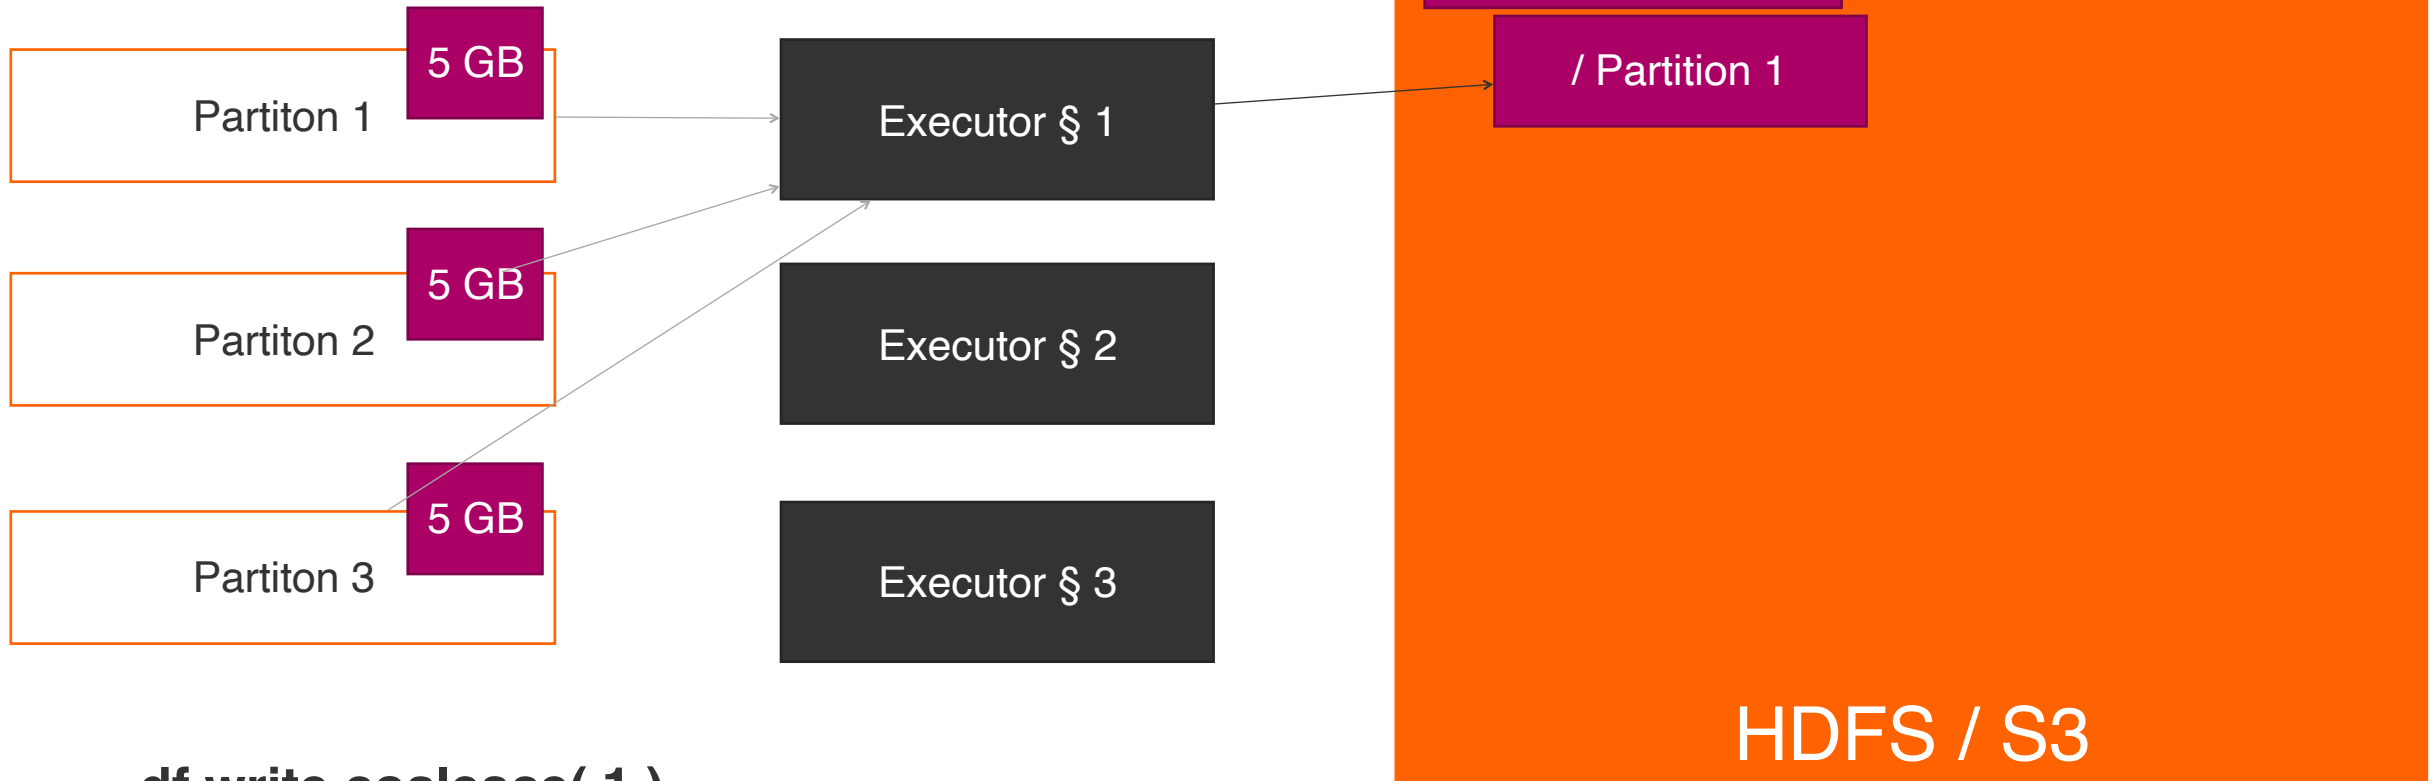
Writing output, cont'd



```
df.write.partitionBy("col1")  
  .parquet("path/output.parquet")
```



Writing output (as one partition)



```
df.write.coalesce( 1 )  
  .parquet("path/output.parquet")
```

Read the partitioned data

```
df.read.parquet("path/Output.parquet")
```

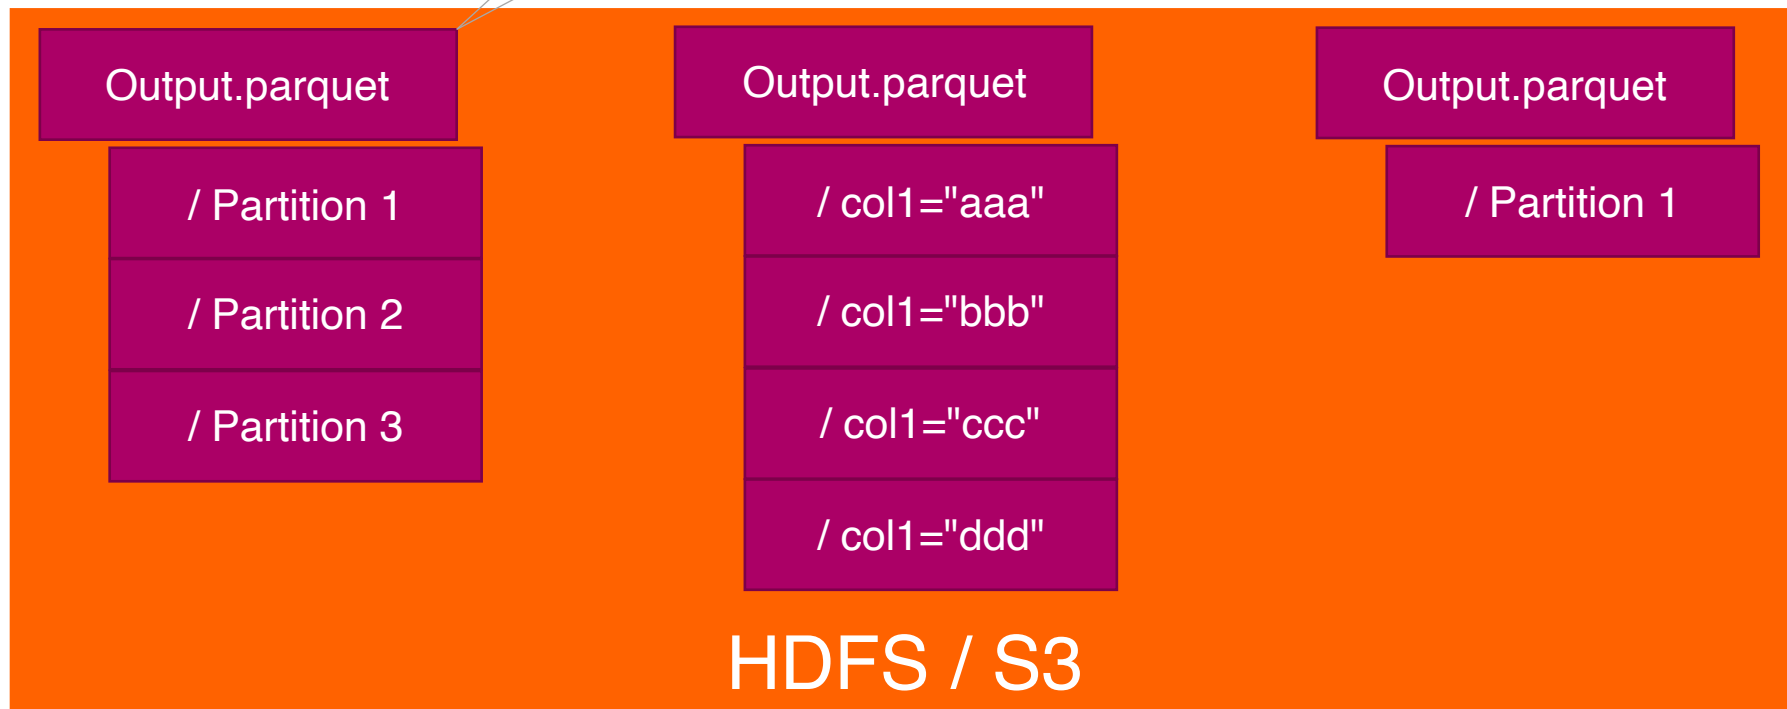
Spark always overrides number of partitions

Partiton 1

Partiton 2

Executor § 1


Executor § 2



Choosing the right data types

	Parquet	CSV	JSON	Avro
Optimised for	File size, SELECT	Readability	Portability	Schema changes
Read speed	Fastest	Slow	Slowest	Fast
Write speed	Slow	Fastest	Slowest	Fast
Partitonable	Yes	Yes	No	Yes
File size	Smallest	Large	Largest	Small
Spark needs to read the whole file	No	Yes	Yes	No

Let's try out the API



<https://github.com/tao-pr/spark-workshop>

Built-in Functions

Column type

```
col("name")  
lit(value)
```



Functions
min()
explode()
rank()
...
collect_set()
udf()



Column type

Built-in Functions

cover	num
Red,Blue,Black	5
Red,Green	10
Green,White	15
Yellow	30



		str	Array[str]
cover	num	colours	
Red,Blue,Black	5	[Red,Blue,Black]	
Red,Green	10	[Red,Green]	
Green,White	15	[Green,White]	
Yellow	30	[Yellow]	

```
dfOut = df.withColumn("colours", split(col("cover"), lit("\\,")))
```


Built-in Functions

product	loss
light	3.0
light	4.0
fridge	8.5
fridge	2.5
fridge	1.5
powersupply	0.1



		double	double
product	loss	logloss	
light	3.0	1.09861229	
light	4.0	1.38629436	
fridge	8.5	2.14006616	
fridge	2.5	0.91629073	
fridge	1.5	0.40546511	
powersupply	0.1	-2.30258509	

```
dfOut = df.withColumn("logloss", log(col("loss"))) )
```

Built-in Functions

product	loss
light	3.0
light	4.0
fridge	8.5
fridge	2.5
fridge	1.5
powersupply	0.1



product	loss	logloss	minlogloss
light	3.0	1.09861229	-2.30258509
light	4.0	1.38629436	-2.30258509
fridge	8.5	2.14006616	-2.30258509
fridge	2.5	0.91629073	-2.30258509
fridge	1.5	0.40546511	-2.30258509
powersupply	0.1	-2.30258509	-2.30258509

```
dfOut = df.withColumn("logloss", log(col("loss")))
        .withColumn("minlogloss", min(col("logloss")))
```

Window functions!

product	loss
light	3.0
light	4.0
fridge	8.5
fridge	2.5
fridge	1.5
powersupply	0.1



product	loss	logloss	minlogloss
light	3.0	1.09861229	1.09861229
light	4.0	1.38629436	1.09861229
fridge	8.5	2.14006616	0.40546511
fridge	2.5	0.91629073	0.40546511
fridge	1.5	0.40546511	0.40546511
powersupply	0.1	-2.30258509	-2.30258509

```
from pyspark.sql.window import Window
```

```
w = Window.partitionBy("product")
```

```
dfOut = df.withColumn("logloss", log(col("loss")) )
```

```
    .withColumn("minlogloss", min(col("logloss")).over(w) )
```

Window functions!

product	loss
light	3.0
light	4.0
fridge	8.5
fridge	2.5
fridge	1.5
powersupply	0.1



product	loss	logloss	rank
light	3.0	1.09861229	1
light	4.0	1.38629436	2
fridge	8.5	2.14006616	3
fridge	2.5	0.91629073	2
fridge	1.5	0.40546511	1
powersupply	0.1	-2.30258509	1

```
from pyspark.sql.window import Window
```

```
w = Window.partitionBy("product").orderBy("loss")
```

```
dfOut = df.withColumn("logloss", log(col("loss")) )
```

```
        .withColumn("rank", rank(col("logloss")).over(w) )
```

Explode array column

product	prices
A	[15,25,35]
B	[]
C	[20]
D	[1,3,5,6]



	Array [double]	double
product	prices	price
A	[15,25,35]	15
A	[15,25,35]	25
A	[15,25,35]	35
C	[20]	20
D	[1,3,5,6]	1
D	[1,3,5,6]	3
D	[1,3,5,6]	5
D	[1,3,5,6]	6

```
dfOut = df.withColumn("price", explode(col("prices"))) )
```

Conditions

product	prices
A	[15,25,35]
B	[]
C	[20]
D	[1,3,5,6]



	double	str
product	price	
A	15	A or B
A	25	A or B
A	35	A or B
B	Null	A or B
C	20	
D	1	
D	3	
D	5	
D	6	

```
dfOut = df.select(col("product"),  
                  explode(col("prices")).alias("price"),  
                  when(col("product").isin(["A","B"]), lit("A or  
B")).otherwise(lit(""))))
```

Using built-in functions anywhere

+	Example
Select	<code>df.select("TARGET", explode(col("a")))</code>
WithColumn	<code>df.withColumn("TARGET", explode(col("a")))</code>
Filter	<code>df.filter(abs(col("v")) > 50)</code>
Aggregation	<code>df.groupBy(floor(col("v"))).agg(sum("w"))</code>

UDF (User-defined function)

Make a function distributable across nodes

```
myfunc = udf( func )
```

`.withColumn("a", myfunc(...))`

`.select(myfunc(...))`

`def func(arg1, arg2)`

`func = lambda arg1, arg2 :`

Spark SQL

Tao Ruangyam, ING Analytics - Frankfurt Hub

Istanbul 2020



Equivalent DataFrame functions on SQL

Interchangeability between **DataFrameAPI** and **SQL command**

```
df.where(col("price")<3000)\  
  .groupBy("grade").agg(  
    count(lit(1)).alias("units"),  
    avg("size").alias("avgsized")  
  ).show(3)
```

grade	units	avgsized
B	113	85.77658681742913
C	122	83.27286338415303
A	16	98.3440670967102

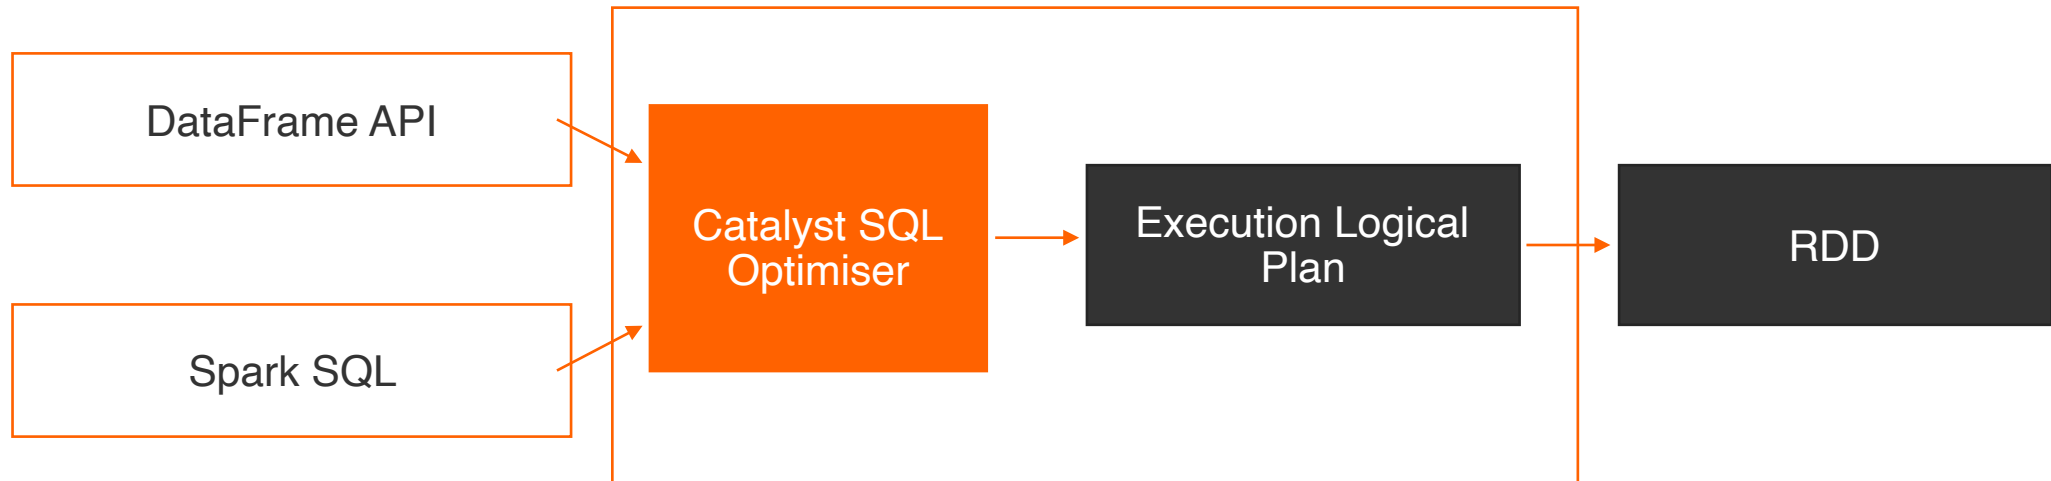
```
spark.sql("""SELECT grade, sum(1) as units, avg(size) as size  
FROM inventory WHERE price<3000  
GROUP BY grade""")  
).show(3)
```

DataFrame API
(SparkSQL)

Equivalent SQL
command

Relation between DataFrame API and SQL engine

Both are compiled into the same form: **Physical execution plan**



Typically users do not have to worry about this box
But the logical plan can also be customisable (Scala)

Instead of writing DataFrame API, we can write SQL

Both **DataFrame API** and **SQL** are compiled into **Physical execution plan**

```
>>> spark.sql("SELECT type, min(price) as low, max(price) as high FROM inventory GROUP BY type").explain()

= Physical Plan =
*(6) HashAggregate(keys=[type#1667], functions=[min(price#1710), max(price#1710)])
+- Exchange hashpartitioning(type#1667, 200)
    +- *(5) HashAggregate(keys=[type#1667], functions=[partial_min(price#1710), partial_max(price#1710)])
        +- *(5) Project [type#1667, price#1710]
            +- *(5) Sort [_nondeterministic#1774 ASC NULLS FIRST], true, 0
                +- Exchange rangepartitioning(_nondeterministic#1774 ASC NULLS FIRST, 200)
                    +- *(4) Project [type#1667, pythonUDF0#5951 AS price#1710, rand(8934628709428066776) AS _nonde
                        +- BatchEvalPython [<lambda>(avgprice#1670L, stdprice#1671L)], [avgprice#1670L, stdprice#16
                            +- *(3) Project [avgprice#1670L, stdprice#1671L, type#1667]
                                +- Generate explode(size#1686), [type#1667, avgprice#1670L, stdprice#1671L], false, [
                                    +- *(2) Project [type#1667, avgprice#1670L, stdprice#1671L, pythonUDF0#5950 AS siz
                                        +- BatchEvalPython [<lambda>(avgsiz#1672L, stdsiz#1673L, qty#1669L)], [avgpri
                                            +- *(1) Project [avgprice#1670L, avgsiz#1672L, qty#1669L, stdprice#1671L, s
                                                +- Scan ExistingRDD[type#1667,grade#1668,qty#1669L,avgprice#1670L,stdpric
```

What is SparkSQL?

Both **DataFrame API** and **SQL** are compiled into **Physical execution plan**

```
>>> df.groupBy("type").agg(min("price").alias("low"), max("price").alias("high")).explain()

= Physical Plan =
*(6) HashAggregate(keys=[type#1667], functions=[min(price#1710), max(price#1710)])
+- Exchange hashpartitioning(type#1667, 200)
    +- *(5) HashAggregate(keys=[type#1667], functions=[partial_min(price#1710), partial_max(price#1710)])
        +- *(5) Project [type#1667, price#1710]
            +- *(5) Sort [_nondeterministic#1774 ASC NULLS FIRST], true, 0
                +- Exchange rangepartitioning(_nondeterministic#1774 ASC NULLS FIRST, 200)
                    +- *(4) Project [type#1667, pythonUDF0#5971 AS price#1710, rand(8934628709428066776) AS _nondeterministic#1774]
                        +- BatchEvalPython [<lambda>(avgprice#1670L, stdprice#1671L)], [avgprice#1670L, stdprice#1671L]
                            +- *(3) Project [avgprice#1670L, stdprice#1671L, type#1667]
                                +- Generate explode(size#1686), [type#1667, avgprice#1670L, stdprice#1671L], false, 1
                                    +- *(2) Project [type#1667, avgprice#1670L, stdprice#1671L, pythonUDF0#5970 AS size#1686]
                                        +- BatchEvalPython [<lambda>(avgsz#1672L, stdsz#1673L, qty#1669L)], [avgsz#1672L, stdsz#1673L, qty#1669L]
                                            +- *(1) Project [avgprice#1670L, avgsz#1672L, qty#1669L, stdprice#1671L, stdsz#1673L]
                                                +- Scan ExistingRDD[type#1667,grade#1668,qty#1669L,avgprice#1670L,stdprice#1671L,avgsz#1672L,stdsz#1673L]
```

Built-in functions in SQL

All Spark built-in functions can also work with SQL!

```
spark.sql("SELECT a,min(b),collect_set(c) FROM table WHERE c>0 GROUP BY a")  
df.where(col("c")>0).groupBy("a").agg(min(col("b")), collect_set(col("c")))
```

Built-in functions in SQL

SQL trick with **explode**



grade	sub
A	[A1, A2, A3]
B	[B1, B2]
C	[C1, C3]
D	[]

```
>>> spark.sql("select grade, explode(sub) from vec").show(10)
```

grade	col
A	A1
A	A2
A	A3
B	B1
B	B2
C	C1
C	C3

Exploded array will **vanish** because explode does not reserve null

Built-in functions in SQL

grade	sub
A	[A1, A2, A3]
B	[B1, B2]
C	[C1, C3]
D	[]

```
>>> spark.sql("select grade, s from vec lateral view outer explode(sub) as s").show(10)
```

grade	s
A	A1
A	A2
A	A3
B	B1
B	B2
C	C1
C	C3
D	null

Reserve null with **lateral view outer**



More complex SQL is also supported

type	grade	seller	price	size	lat	lng
Apartment	A	24	4262.7754	92.717896	53.396652	35.058865
Apartment	B	29	2935.234	71.07544	56.5649	41.392433
WG	B	6	2136.9932	111.69808	24.81349	14.496414
WG	C	17	988.5911	43.284058	24.123186	38.55362
...						
Apartment	C	14	1091.334	30.537722	17.6535	30.582376

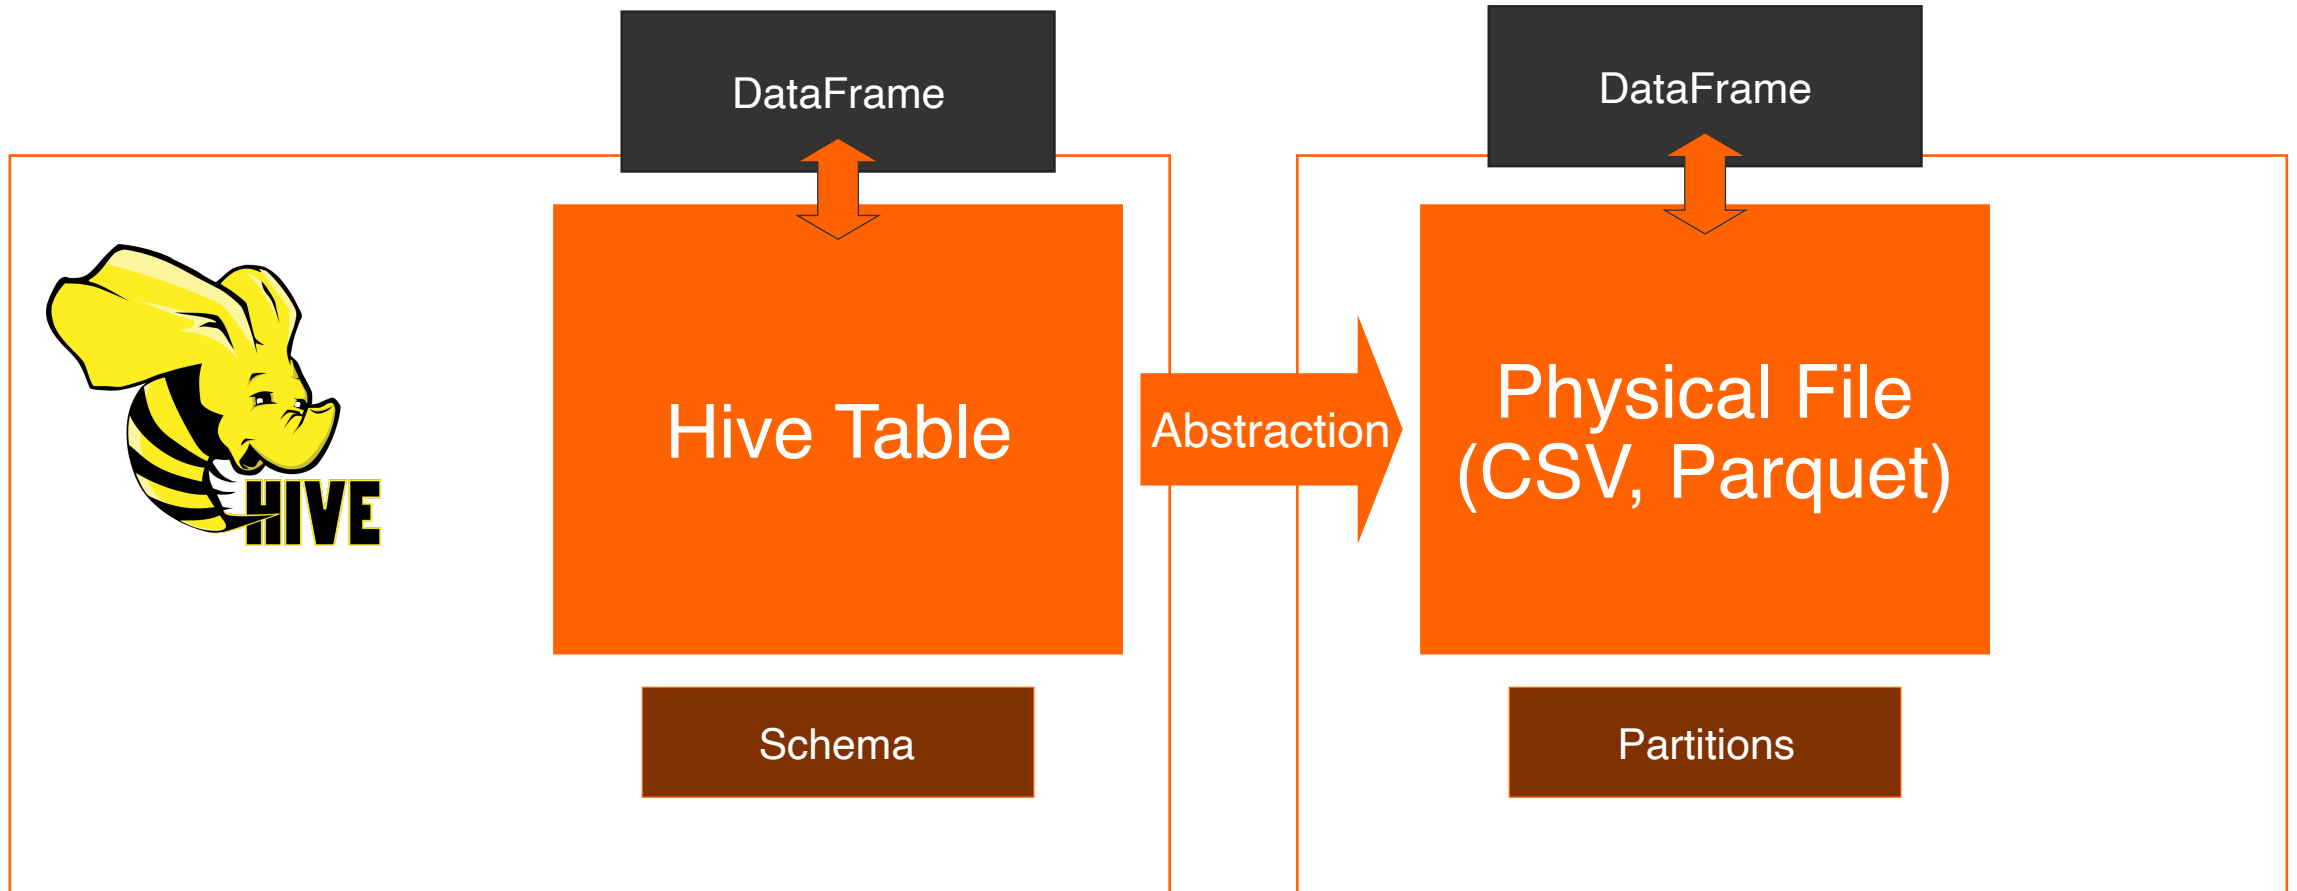
```
SELECT type, grade, price, COALESCE(
  CASE WHEN size<30 THEN "small" ELSE NULL END,
  CASE WHEN size<80 THEN "medium" ELSE NULL END,
  CASE WHEN size<100 THEN "large" ELSE "xlarge" END) AS size from df
```

type	grade	price	size
Apartment	A	3271.7893	large
Apartment	B	3043.0178	medium
WG	B	2311.949	xlarge
WG	C	832.8406	medium
Apartment	B	3210.3003	medium
Apartment	B	3756.9866	large
...			
House	A	8788.338	xlarge
House	C	3782.6099	xlarge
WG	B	1800.9767	large
Apartment	C	790.08777	small

You can use SQL at any complexity
Spark SQL can handle like other industry-grade RDBMS

SparkSQL supports Hive

Hive is a **relational database** engine for Hadoop



SparkSQL supports Hive

Hive table abstracts the access to physical files with SQL

```
df = spark.read.csv("path/file.csv").filter(col("a") == 100).select("a","b")

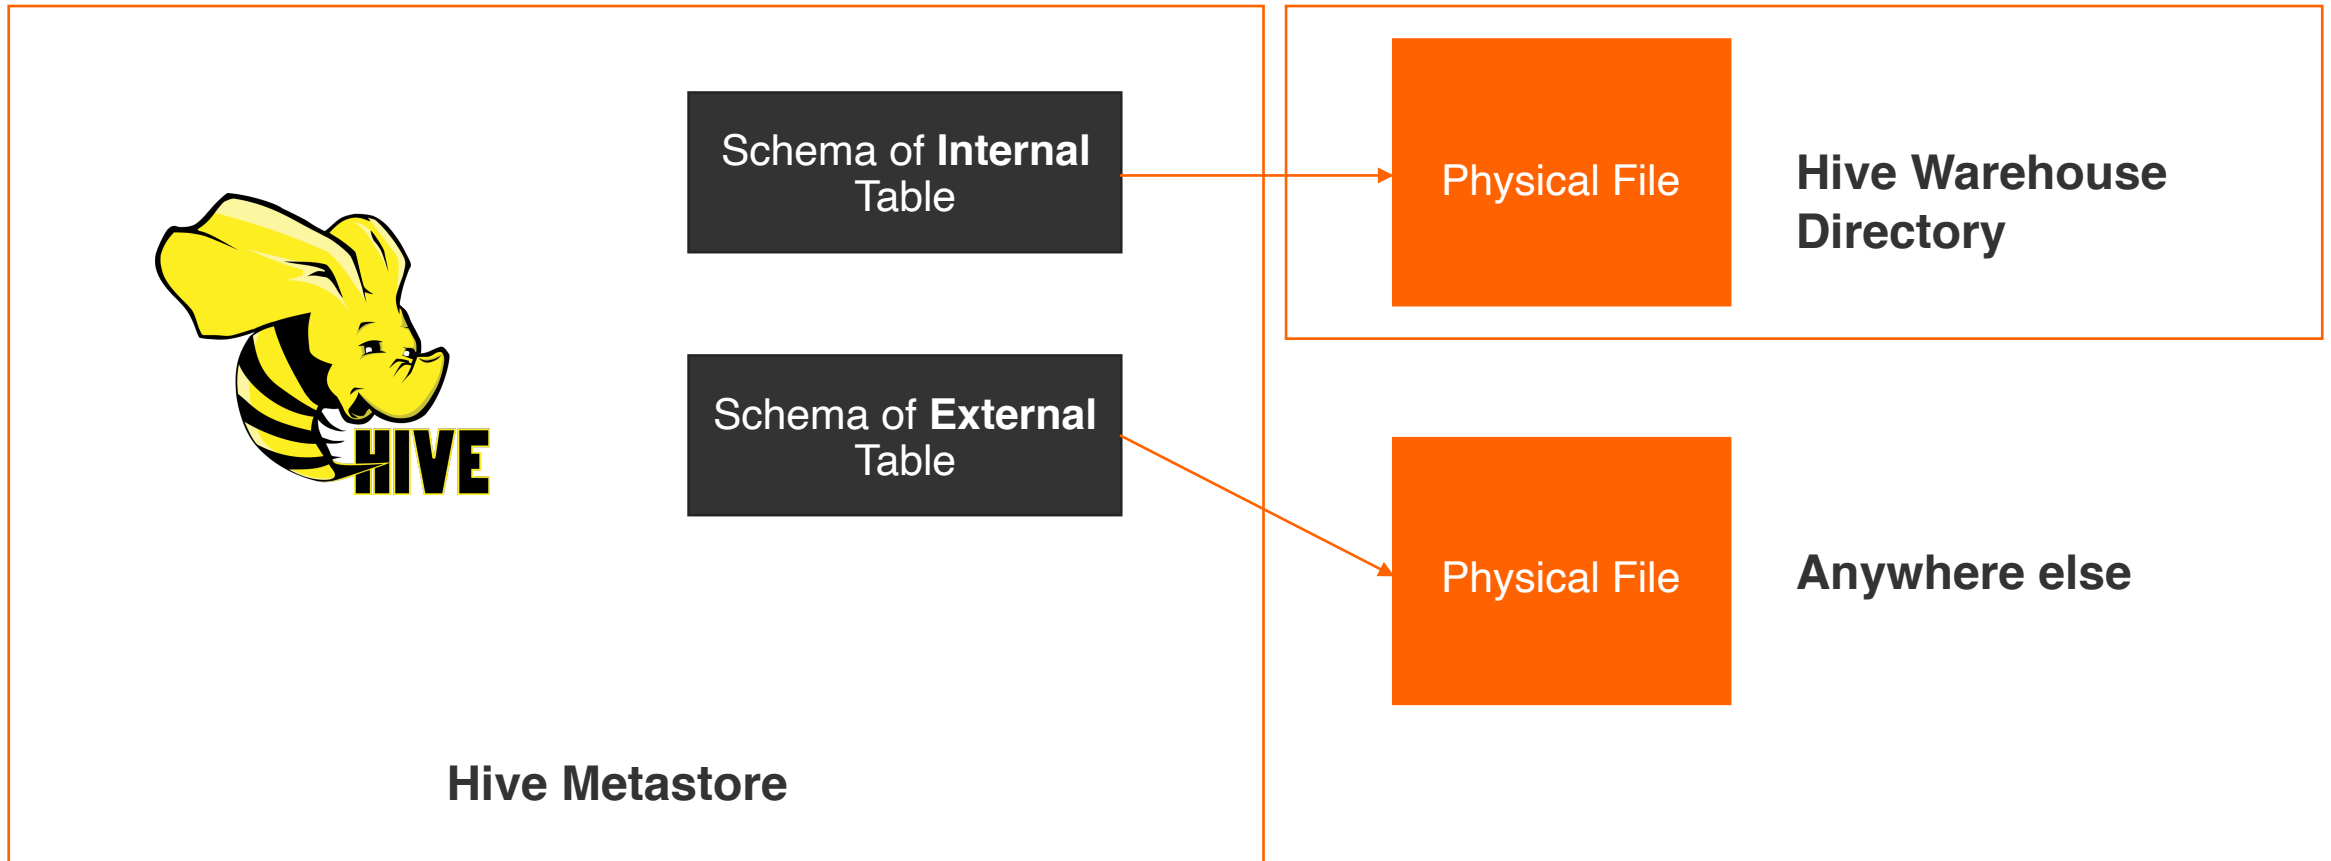
df = spark.sql("SELECT a, b FROM schema.table WHERE a = 100")
```

To create a Hive table from an existing file

```
CREATE EXTERNAL TABLE IF NOT EXISTS schema.table(a INT, b STRING)
ROW FORMAT DELIMITED
FIELD TERMINATED BY ','
STORED AS CSV
LOCATION '/path/file.csv'
```

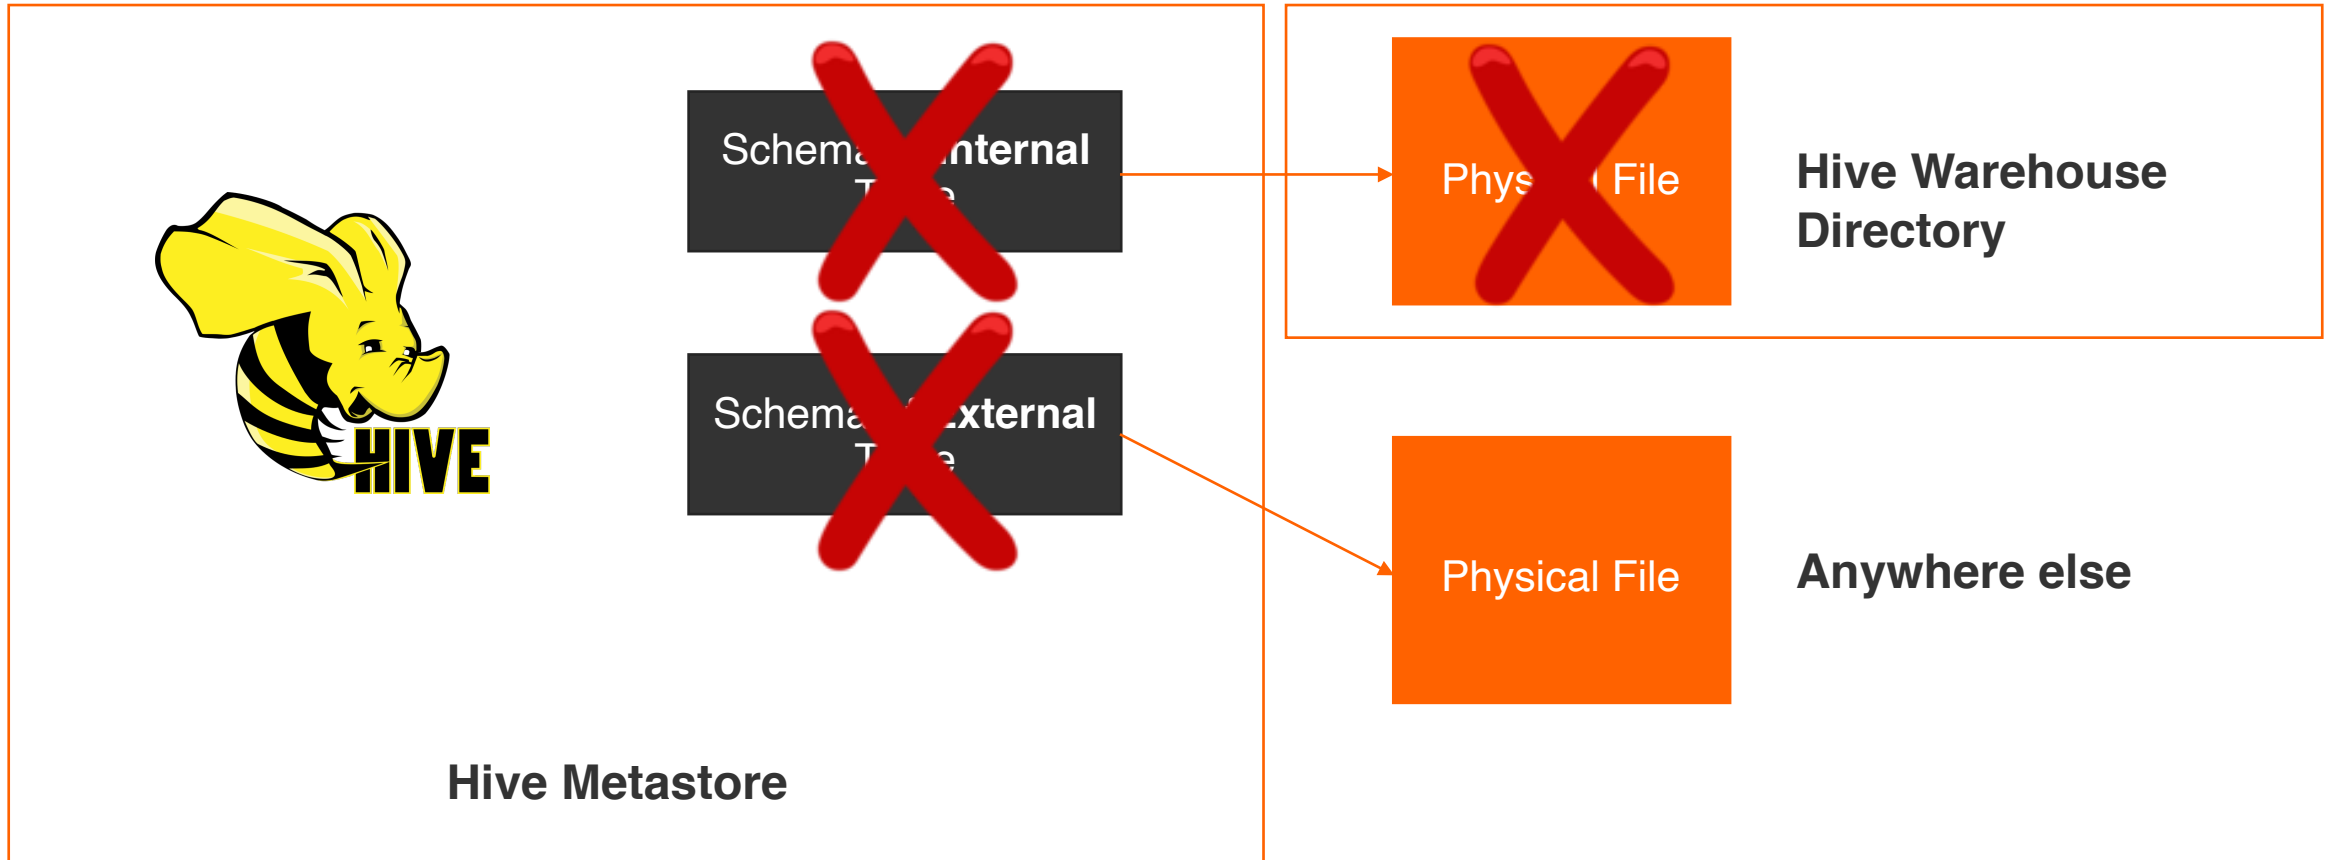
Hive External vs Internal table

External means outside of Hive's Warehouse location



Hive External vs Internal table

Dropping external table also destroys the physical file



Hive vs Impala

Impala is upto 70x faster as it cache the direct pointer to the physical file

