# Spark Workshop

Tao Ruangyam, ING Analytics – Frankfurt Hub

Istanbul 2020

ING

# Spark Components

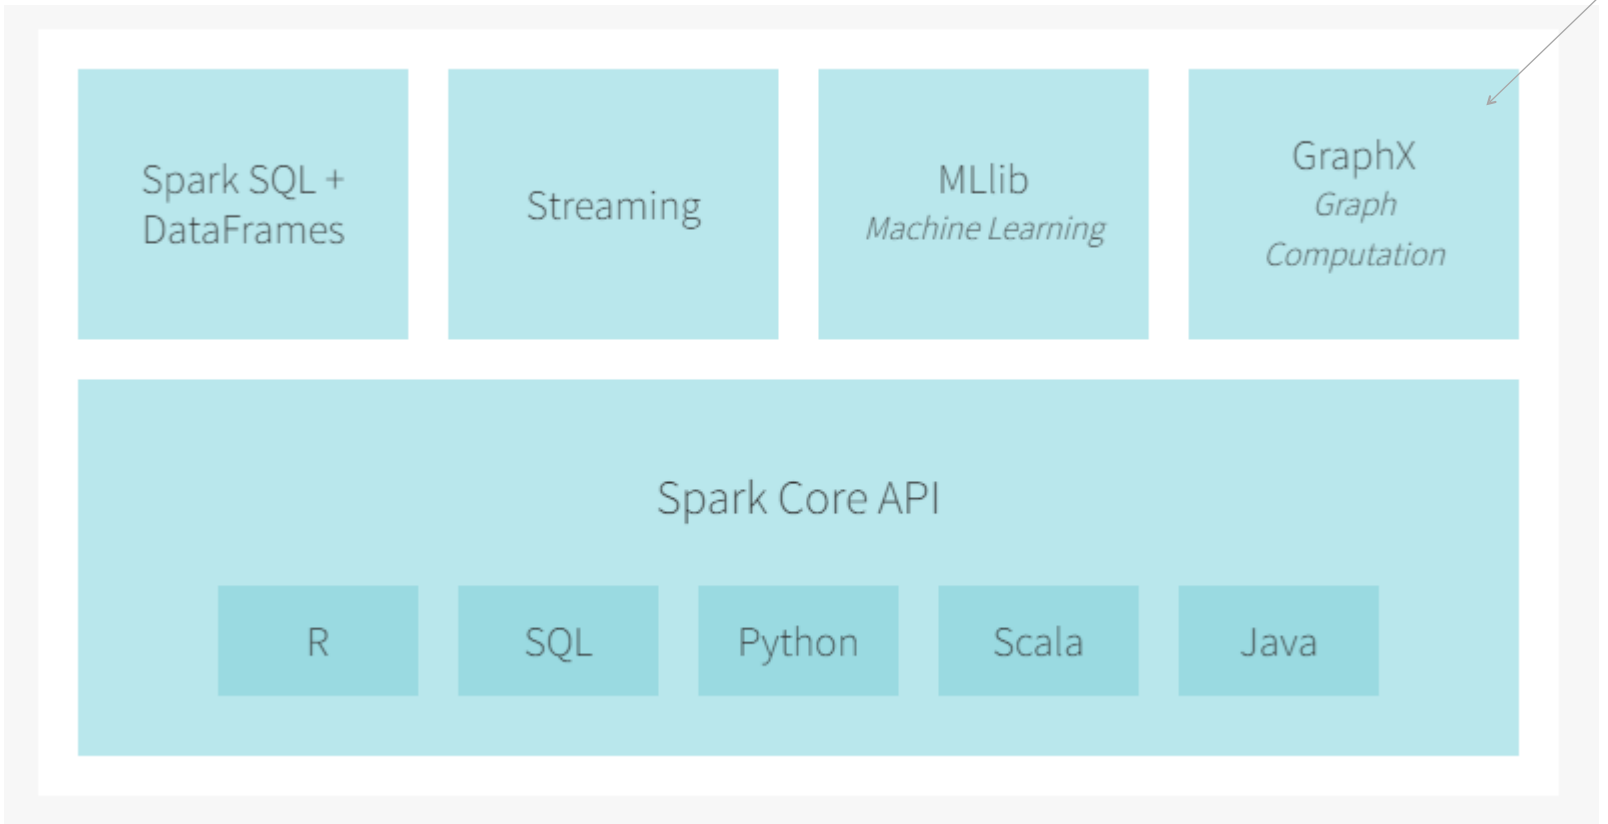


Spark 3.0 is aimed to have **Neo4J's Cypher** in built.

Spark 2.x supports up to JDK8, But Spark 3 will support JDK11

# This workshop covers


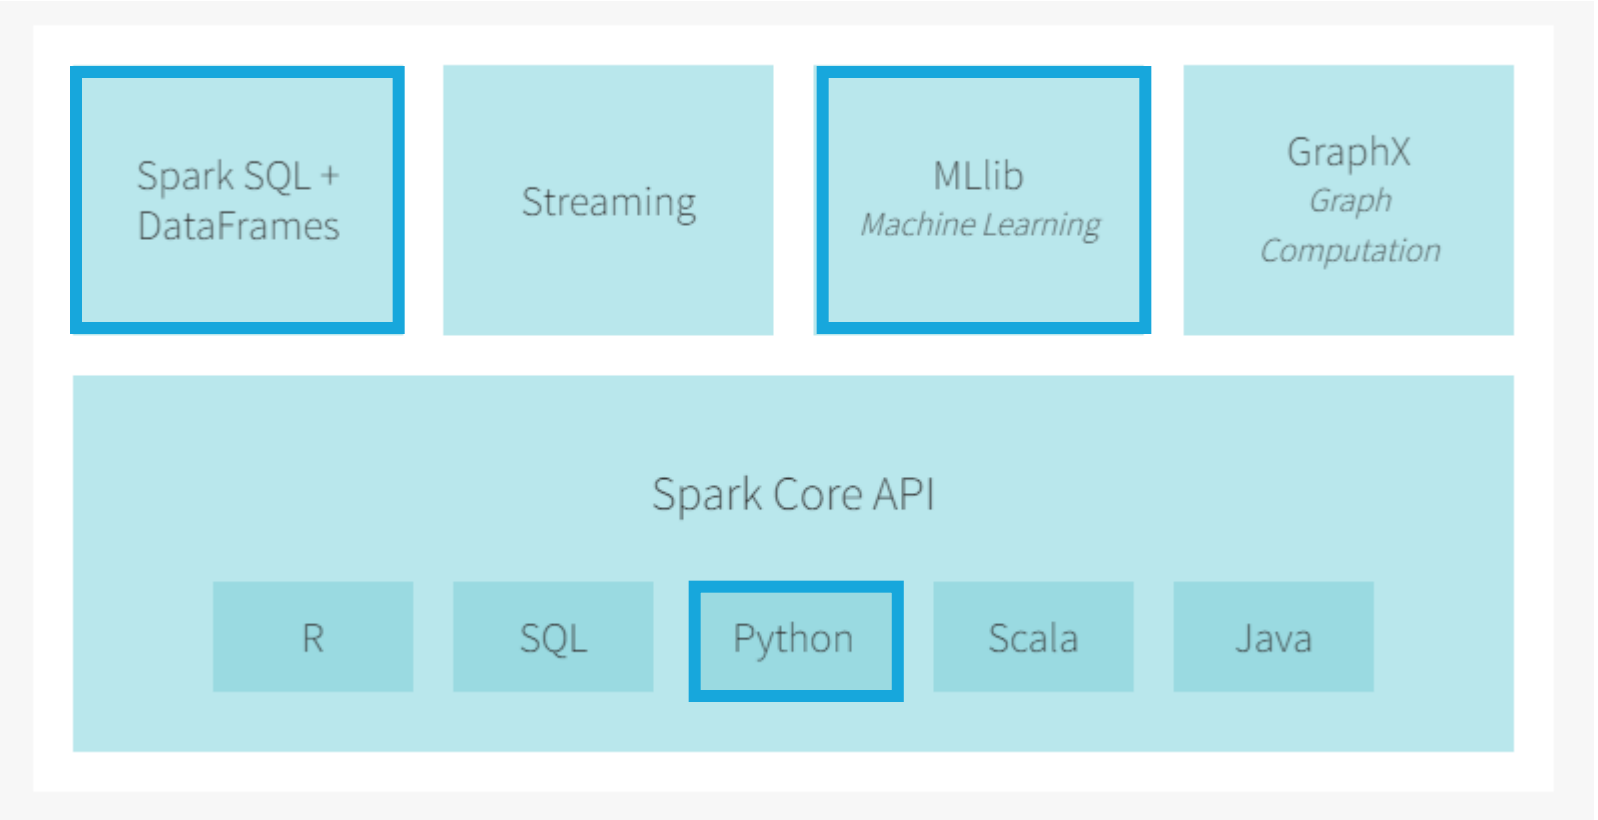




ING

# Agenda

| Day 1 | | Coding? |
|---|---|---|
| 10:00 – 10:50 | Intro to Spark Ecosystem | - |
| 11:00 – 12:00 | Spark APIs | - |
| Lunch break | | |
| 13:00 – 13:50 | Spark APIs (cont) | Yes |
| 14:00 – 15:00 | Spark Memory management & Optimisation | - |

| Day 2 | | Coding? |
|---|---|---|
| 10:00 – 10:50 | Spark SQL | Yes |
| 11:00 – 12:00 | Spark ML | - |
| Lunch break | | |
| 13:00 – 13:50 | Spark ML (cont) | Yes |
| 14:00 – 15:00 | Wrap up, Best practices & Tips | Optional |

# Spark Ecosystem

Spark capabilities

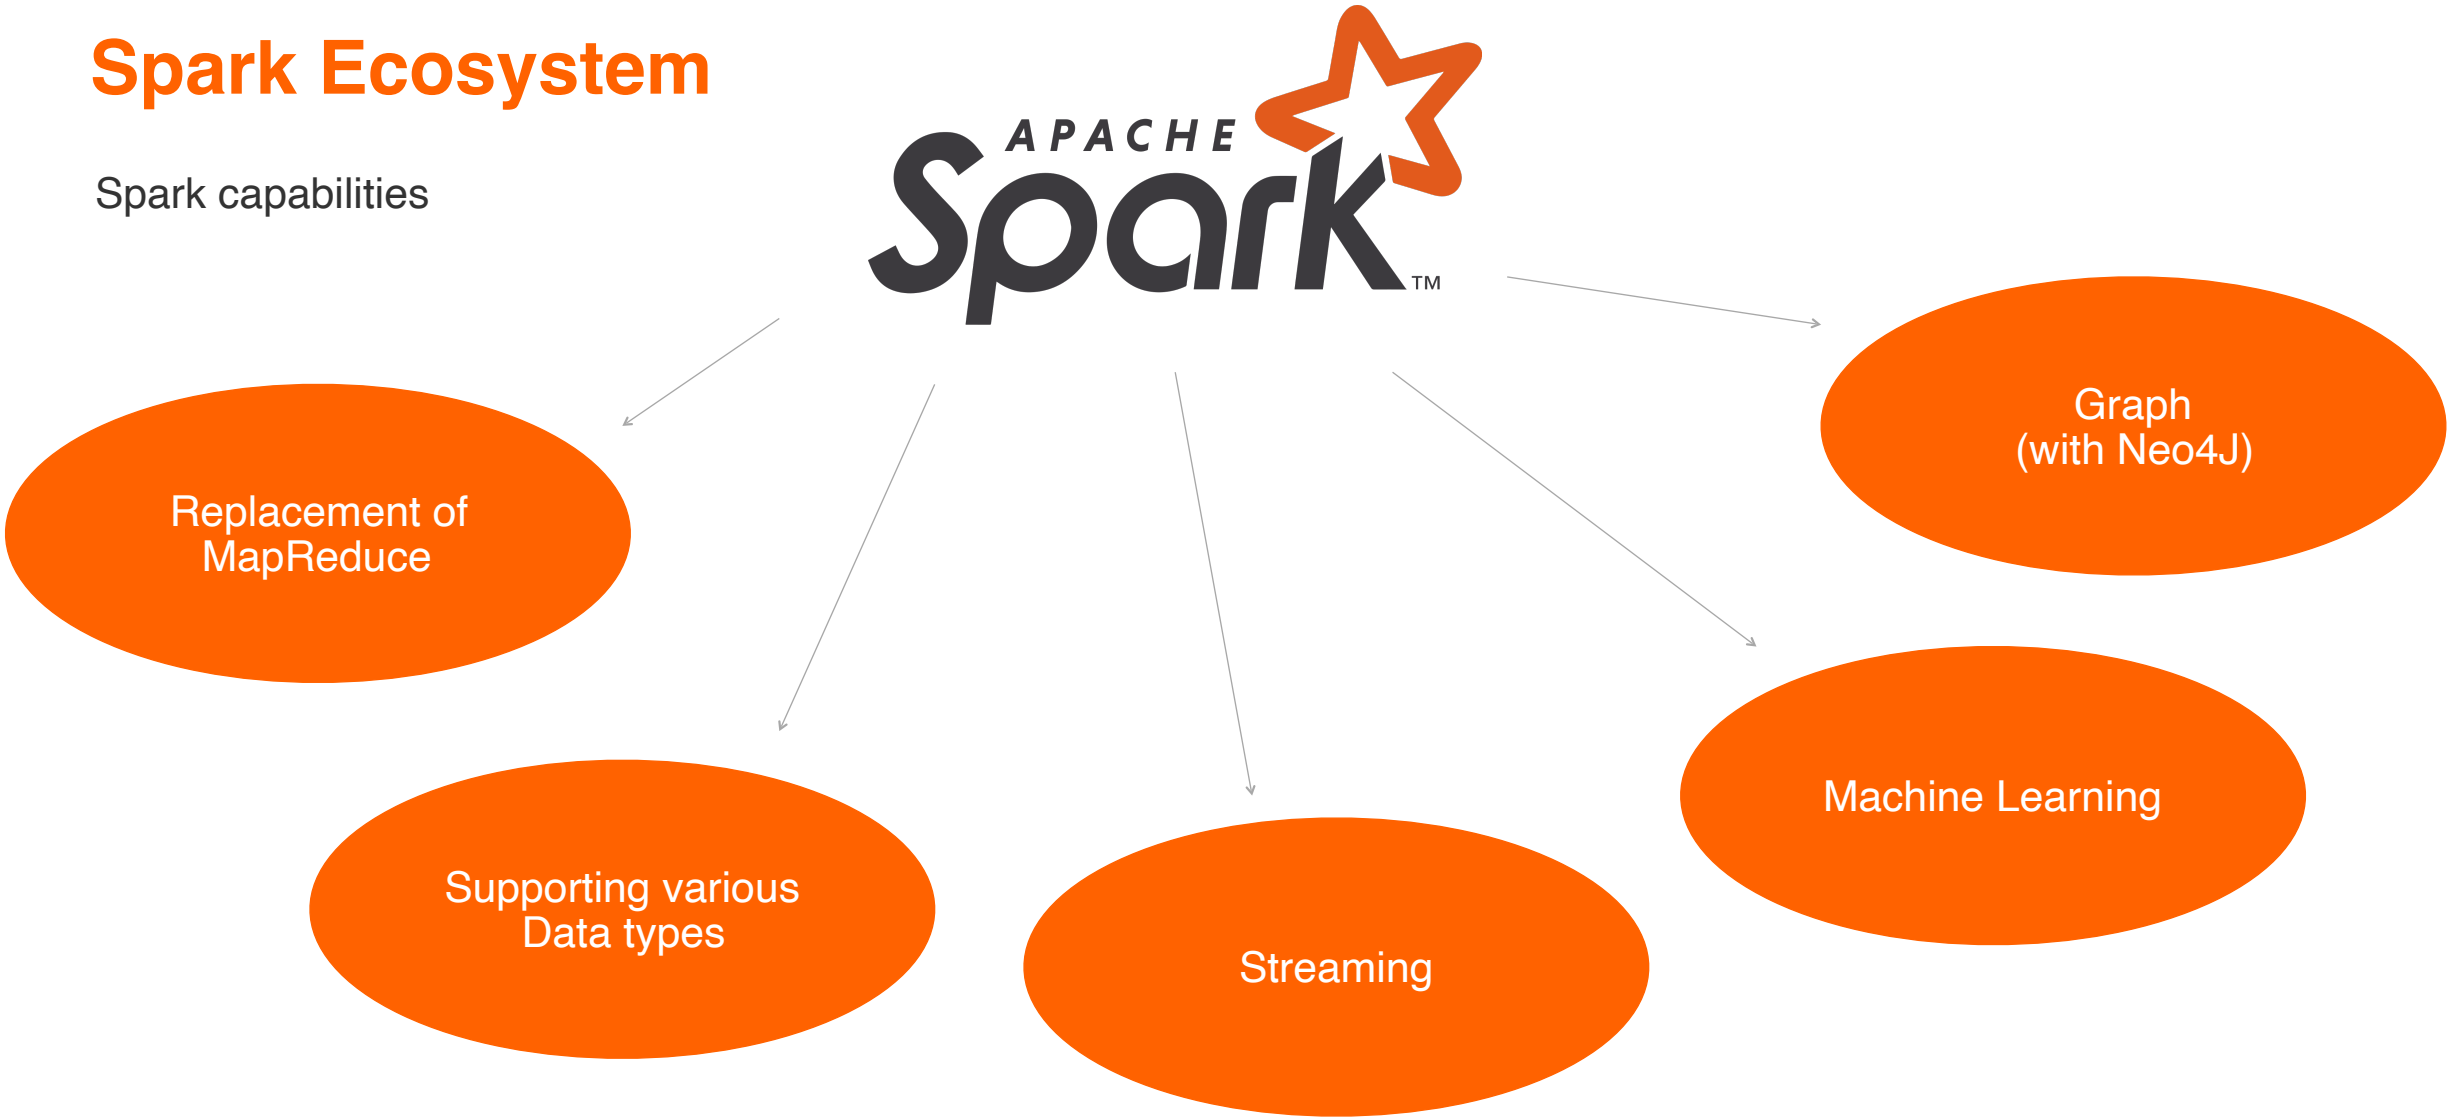




ING

# Spark vs Single-Machine

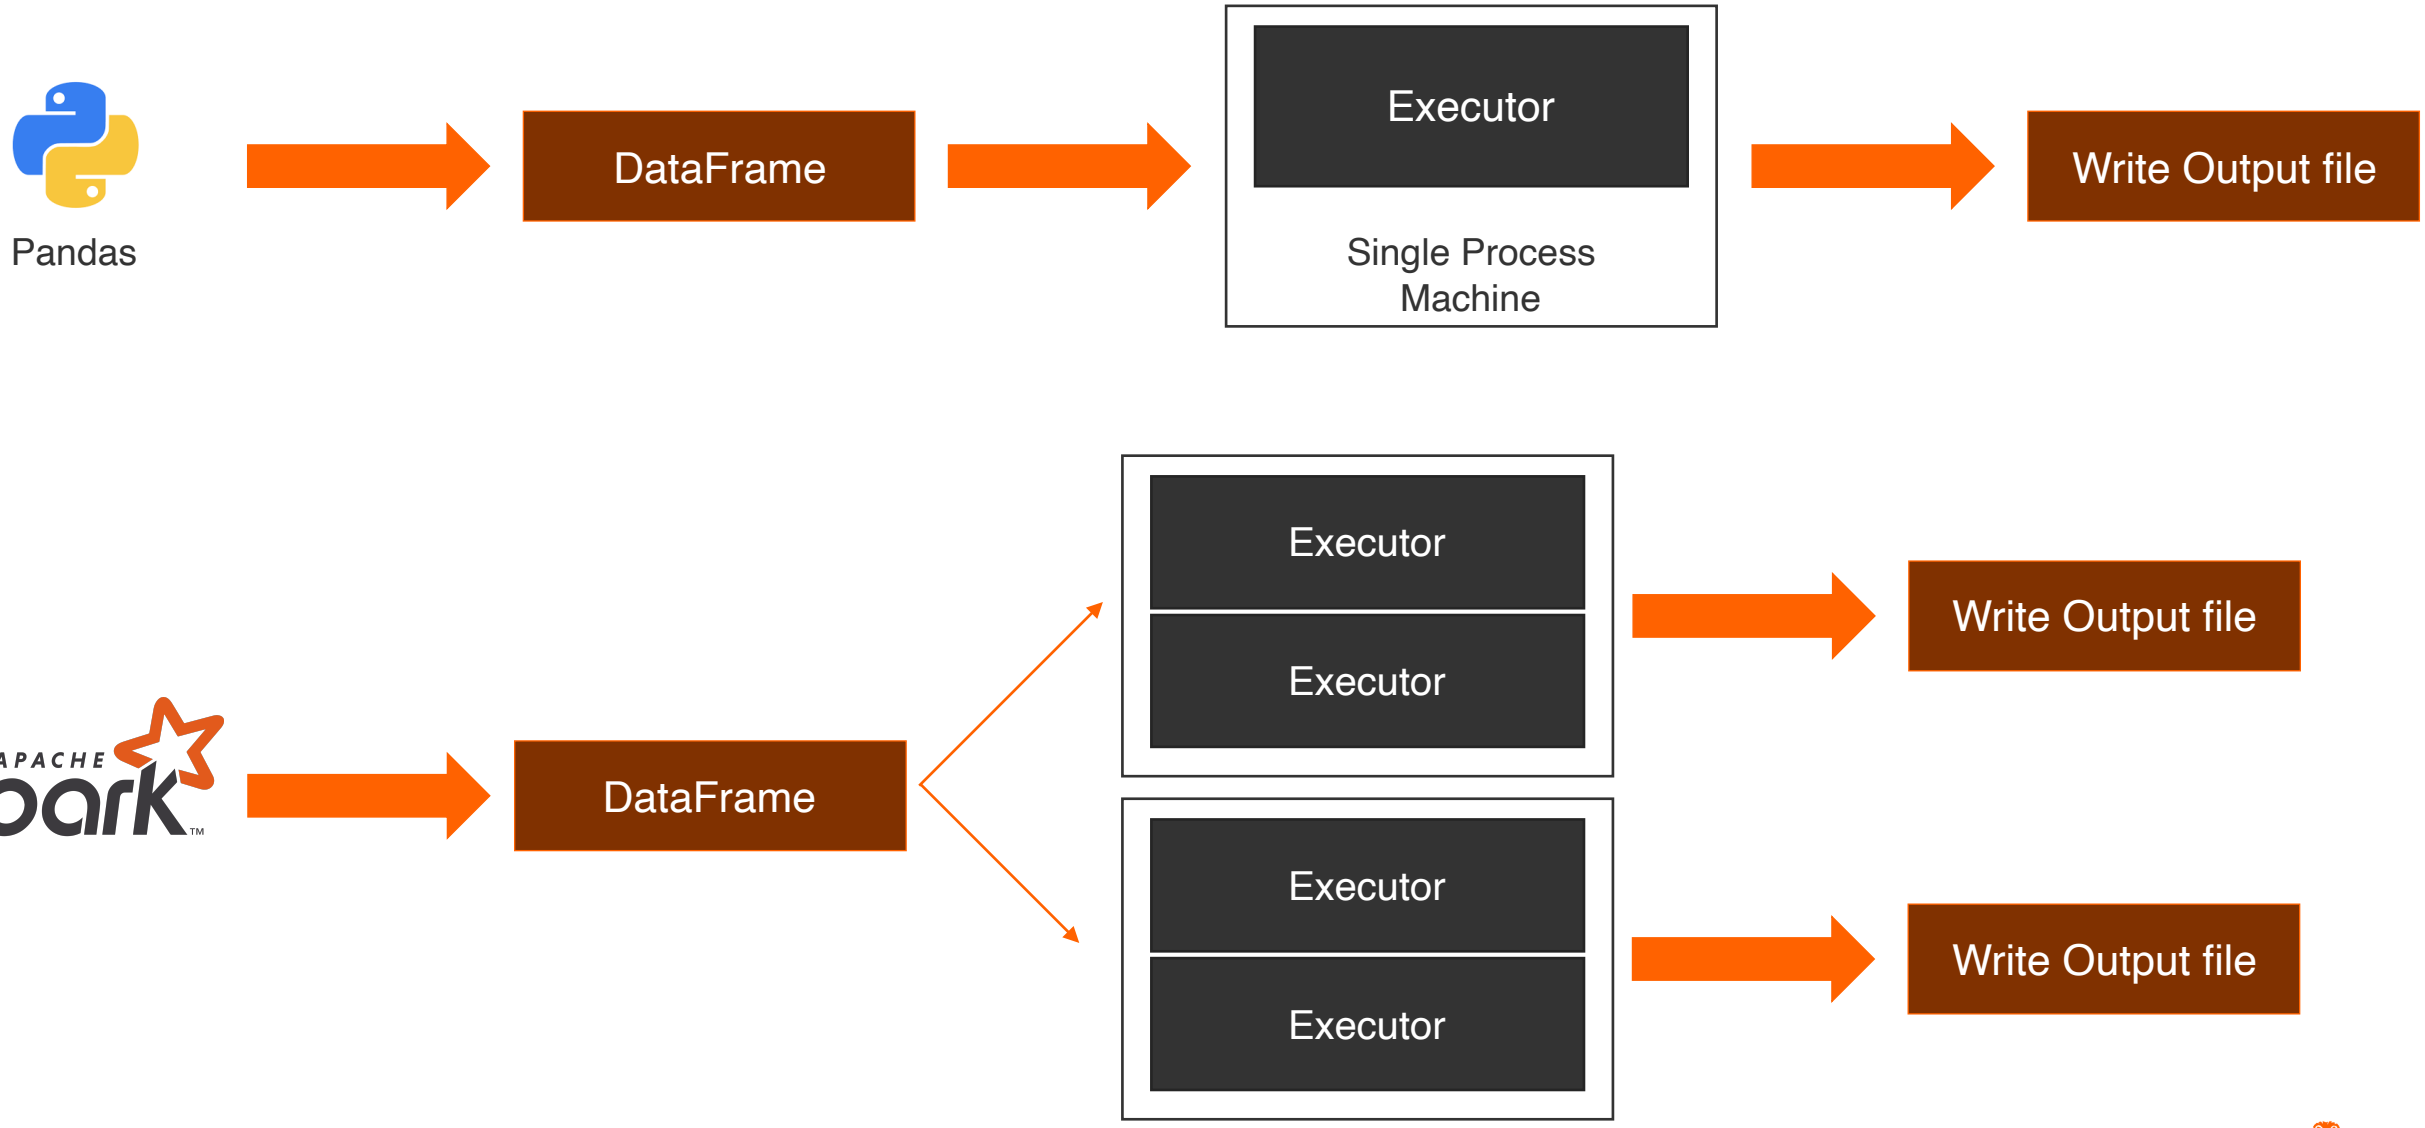




ING

# Spark Environment

Orchrestration model of Spark cluster

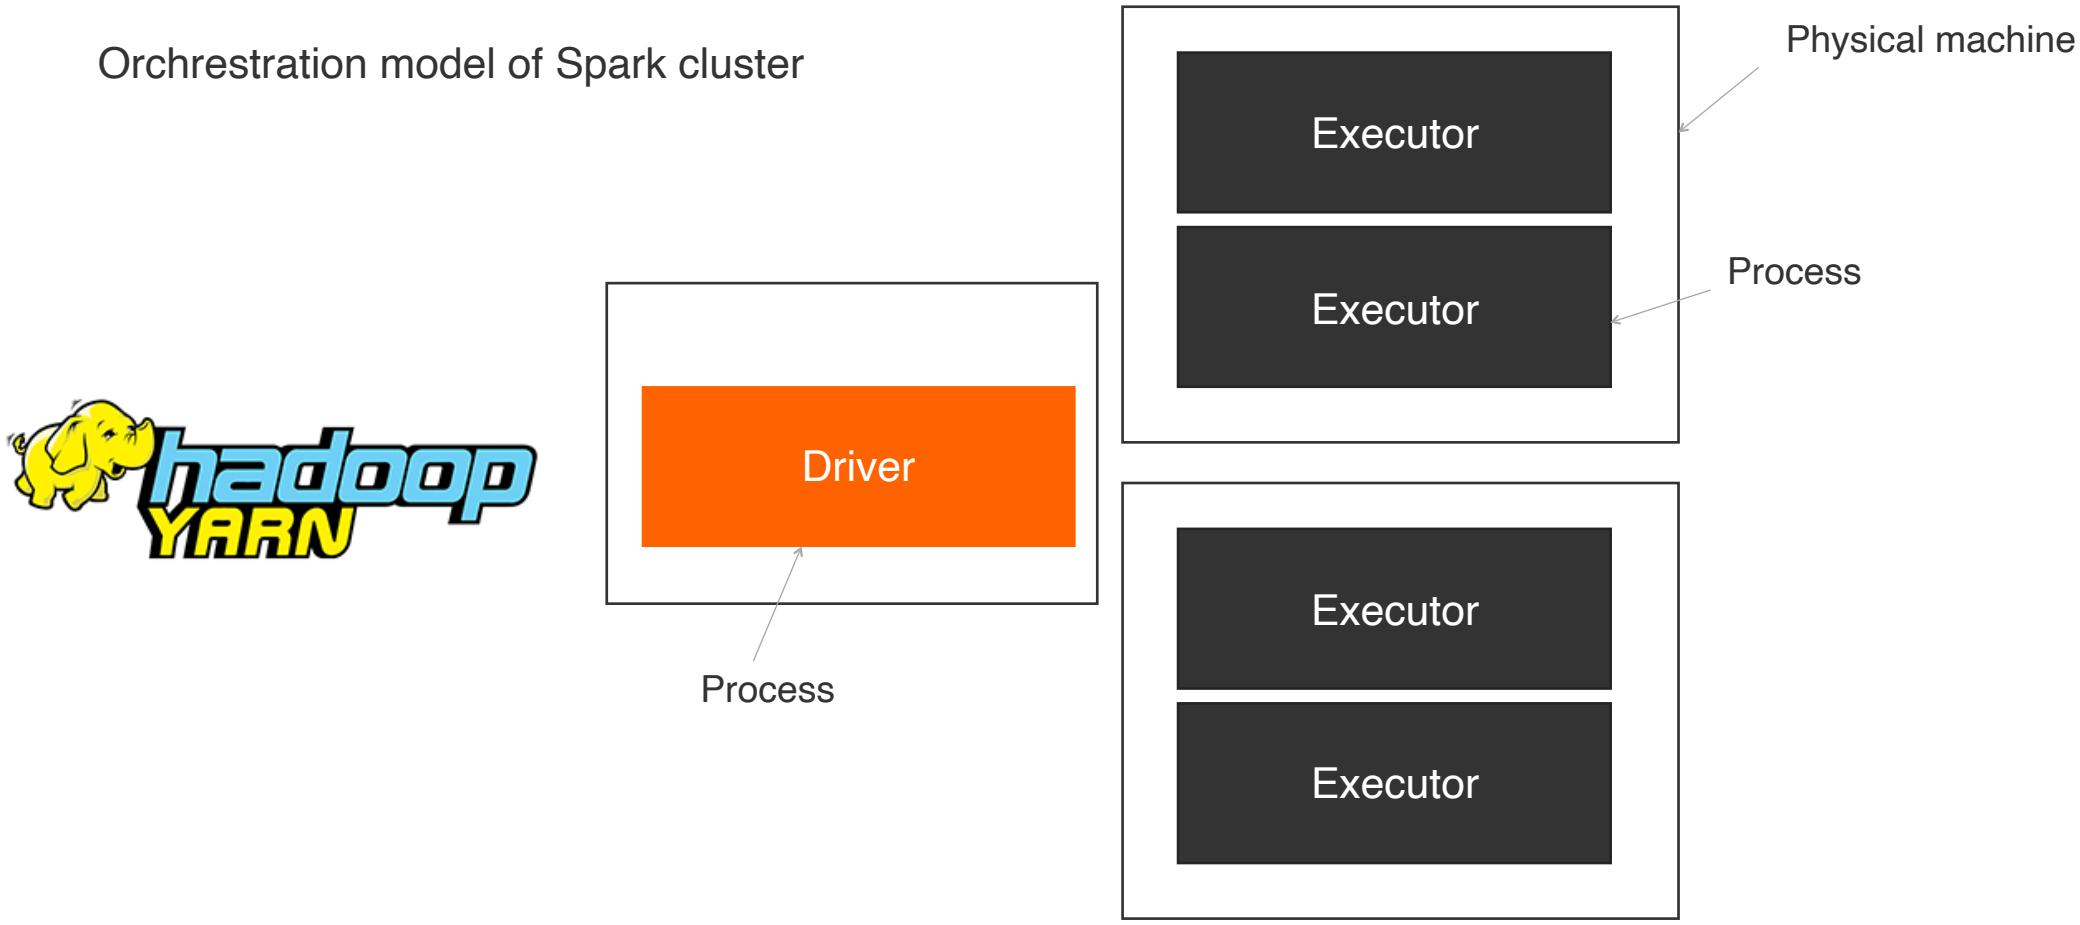

# Spark Environment

Orchrestration model of Spark cluster


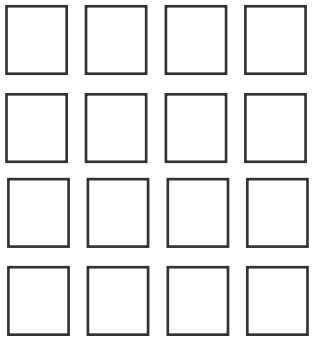

Machines on
Cluster



ING

# Spark Environment

Orchrestration model of Spark cluster


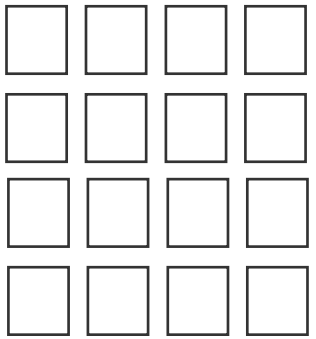

User submits a Spark job

Machines on Cluster




ING

# Spark Environment

Orchrestration model of Spark cluster

# Spark Environment

Orchrestration model of Spark cluster


Communicate with Driver

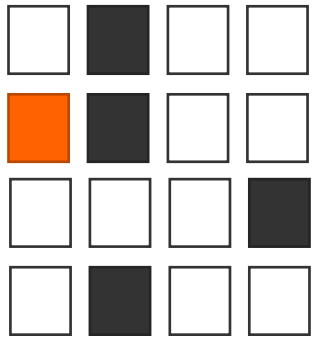

Machines on Cluster

Driver assigns executors, define tasks and distribute them to executors

# Spark Environment

Orchrestration model of Spark cluster

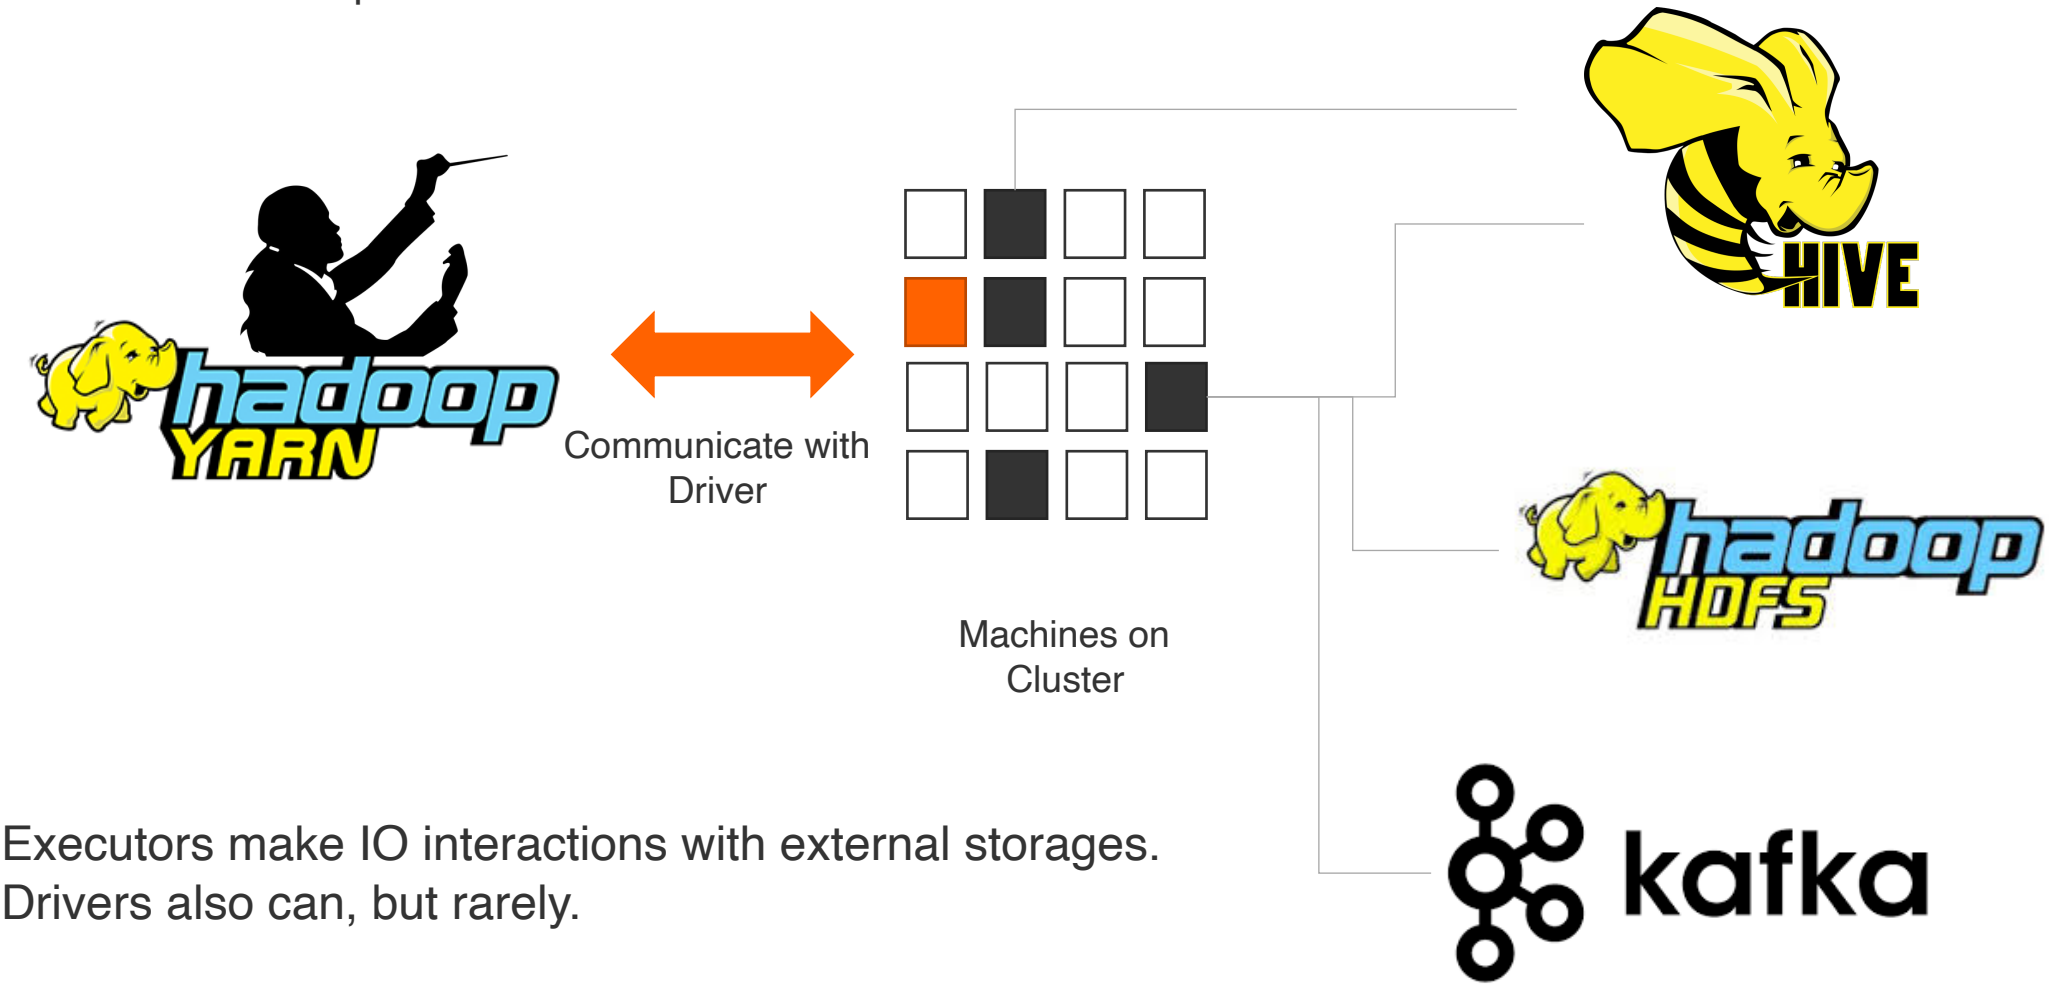


Executors make IO interactions with external storages.
Drivers also can, but rarely.



ING

# Spark Environment

Orchrestration model of Spark cluster

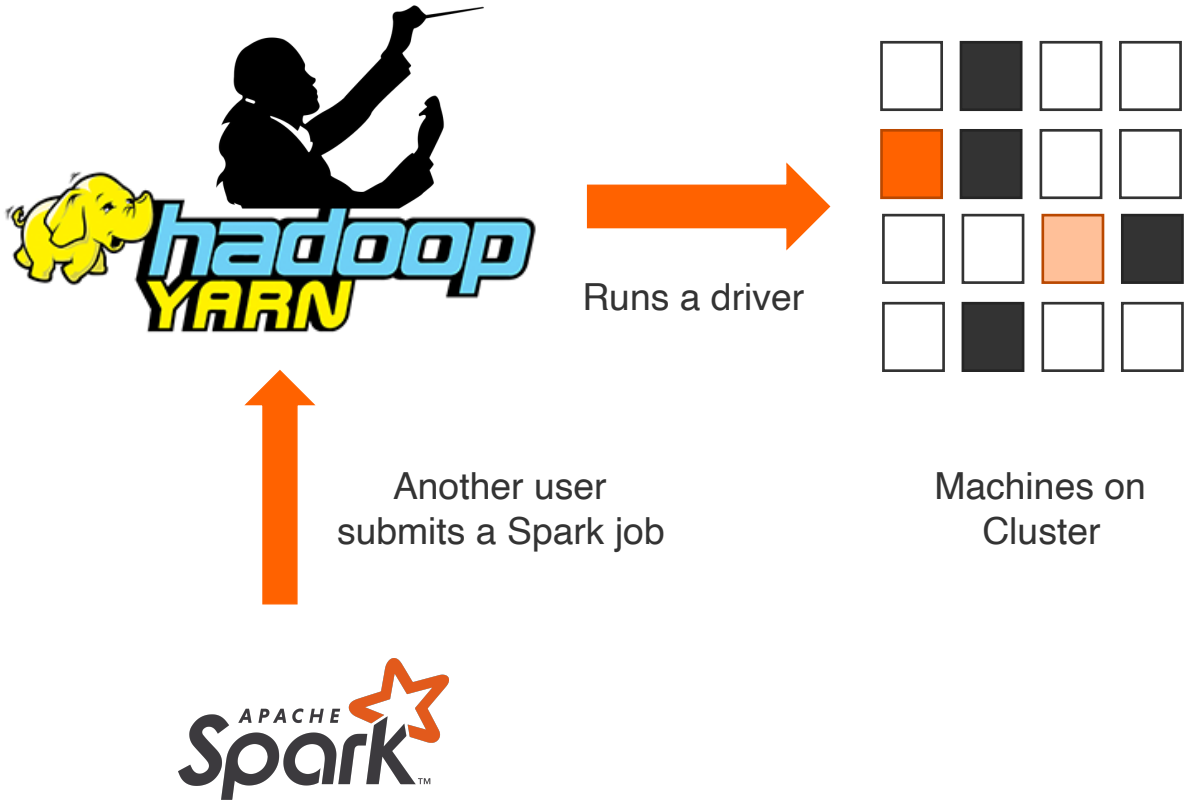




ING

# Spark Environment

Orchrestration model of Spark cluster

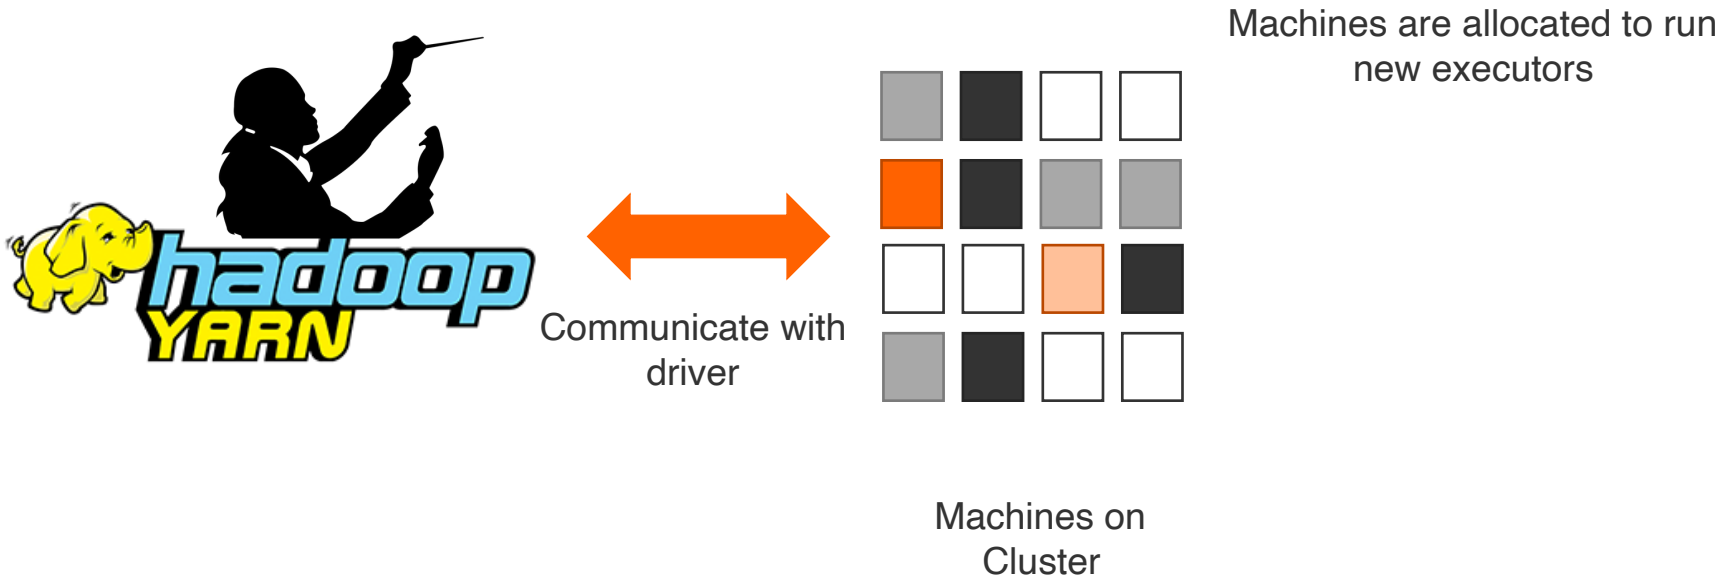


**NOTE:** One physical machine can run more than one executors.

# Spark Environment

Orchrestration model of Spark cluster

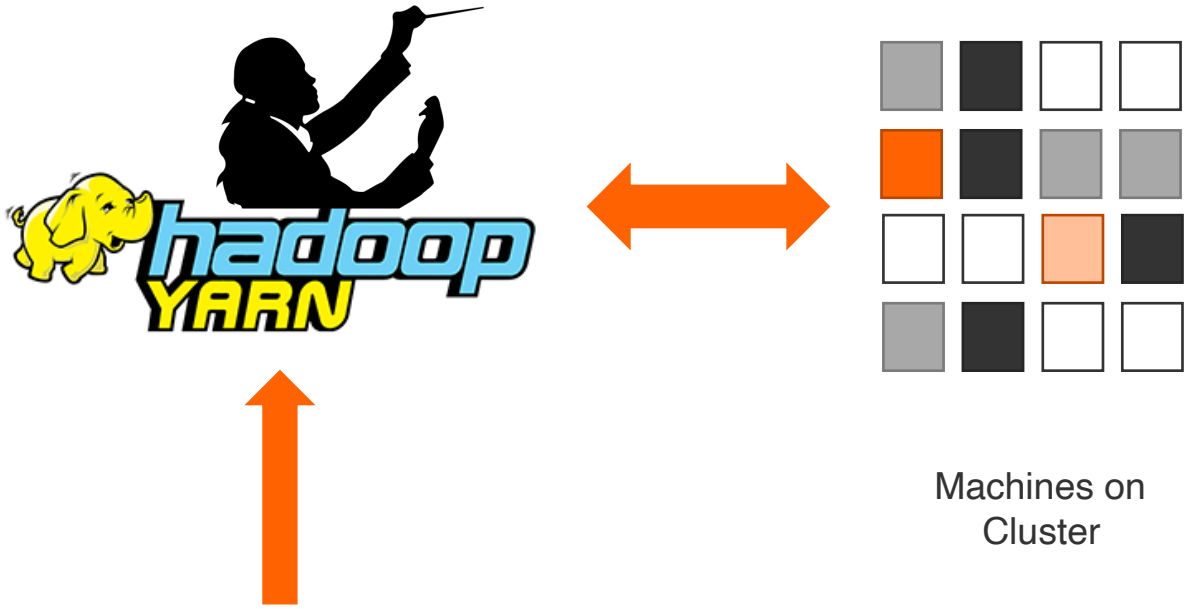


Machines on Cluster

1. Driver takes forever to start
2. Though the driver starts, during the run the job halts.

Large job submitted, requiring lots of resources beyond the supply.



ING

# Contributors and Providers of Spark

Official merger began in 2019




| Contributor of Spark codebase | Vendor of Spark distribution **(CDH)** | Vendor of Spark distribution **(HDP)** |
|---|---|---|
| MLLib | Parquet (collab. with **Twitter**) | ORC (collab. with **Facebook**) |
| DeltaLake | Hue | NiFi |
| MLFlow | Impala | |



ING

# Spark is 72% written in Scala

| | With Scala | With Python |
|---|---|---|
| **Performance** | 10x faster,<br>more memory efficient | |
| **Object serialisation** | All types, case classes are natively supported | Native python types, not with numpy types |
| **RDD/Dataframe API** | Y | Y |
| **Typed Dataset API** | Y | N |
| **Notebook** | Yes, Zeppelin | Yes, Jupyter |
| **Types error** | Compilation time | Runtime only |

# PySpark vs Scala Spark

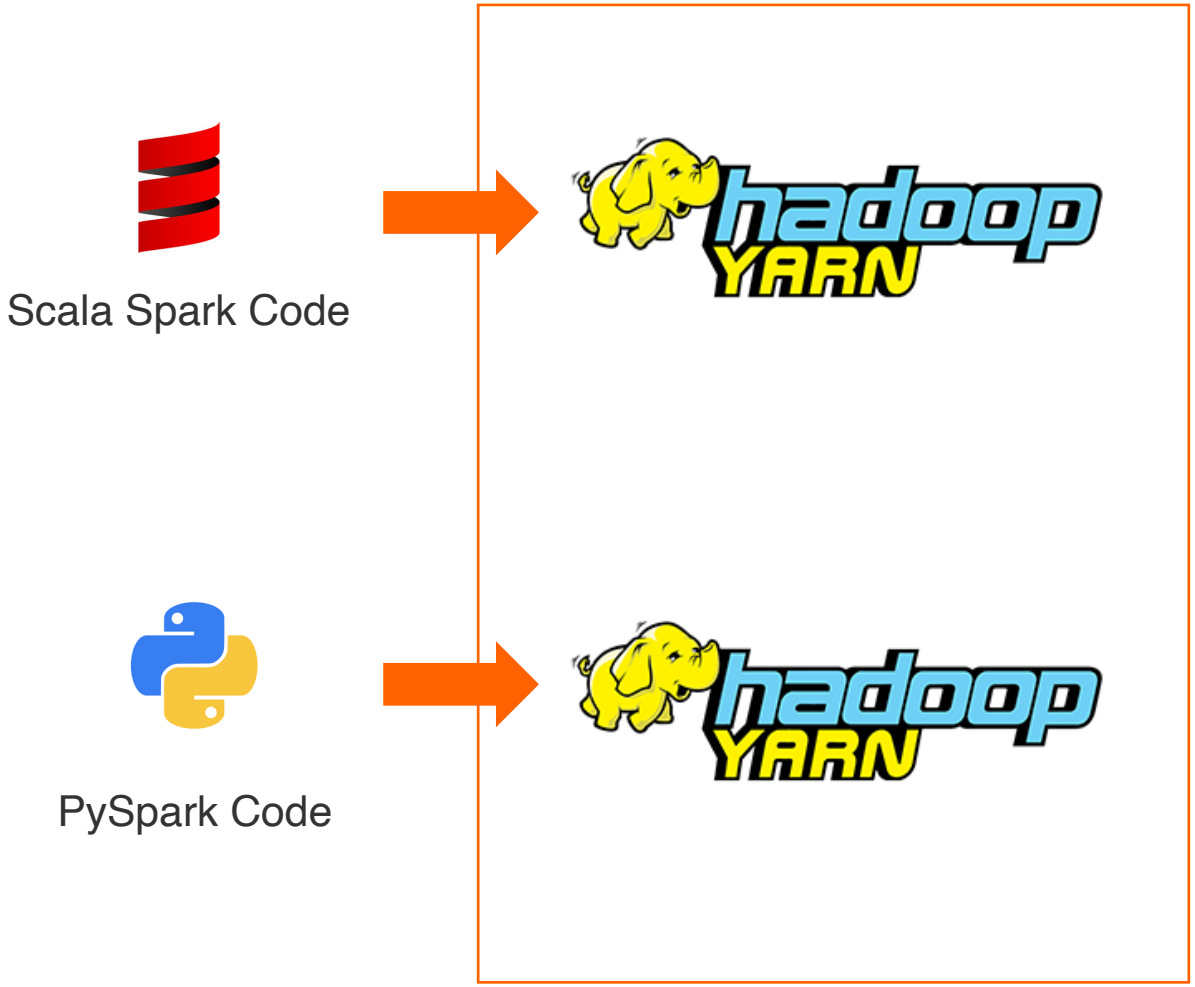


**Executor Memory Region**

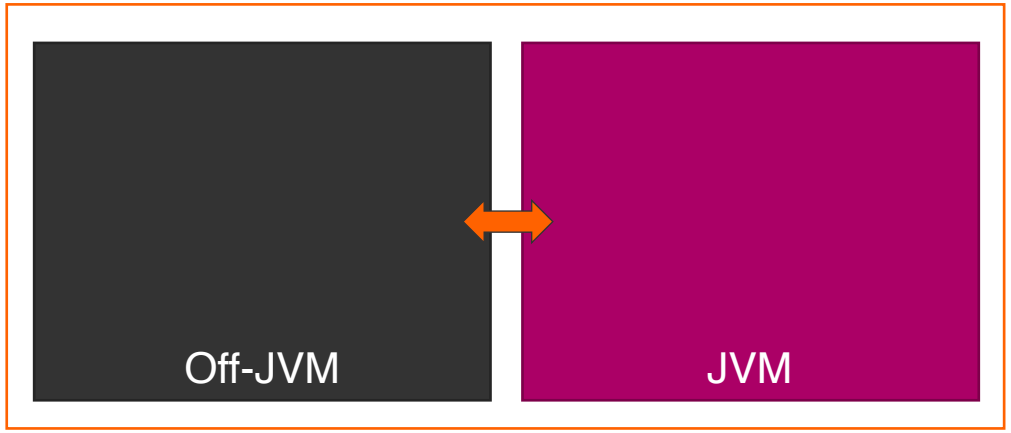


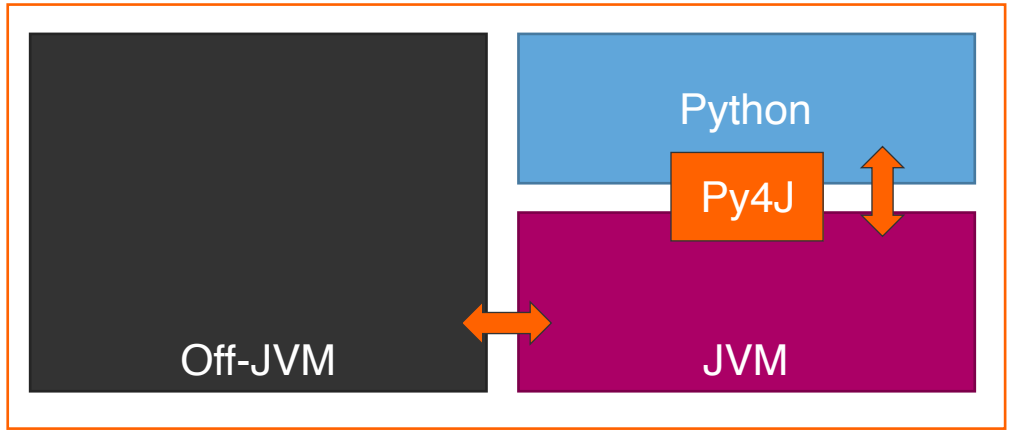




ING